



**ifis**

Institut für Informationssysteme  
Technische Universität Braunschweig

# **Relational Database Systems I**

**Wolf-Tilo Balke**

**Niklas Kiehne, Denis Nagel,  
Enrique Pinto Dominguez**

Institut für Informationssysteme  
Technische Universität Braunschweig  
<http://www.ifis.cs.tu-bs.de>



# 0. People involved

- Who is who ?
  - Wolf-Tilo Balke
    - lecture, exams
  - Niklas Kiehne
    - lecture, detours
  - Denis Nagel, Enrique Pinto Dominguez
    - homework / tutorials
  - Regine Dalkiran
    - office
- In case of questions,  
feel free to ask us!





# 0. Organizational Issues

- Lecture
  - October 23, 2025 to February 5, 2026
  - 15:00 – 17:15 (including a break)
  - *integrated lecture*  
(theory and detours)
  - 5 credits
  - Written exam on **13.03.2026**
- Homework
  - weekly assignments
    - ... can be downloaded from StudIP
    - ... must (!) be completed in groups of **three** students





# 0. Organizational Issues

- **Tutorial groups**
  - led by our HiWis
  - homework discussion
- In order to pass this **module** (5CP) you need to
  - I) ... achieve **50%** of homework points  
(Studienleistung, ungraded)
  - 2) ... pass the exam (Prüfungsleistung, graded)





# 0. Homework and Tutorials

- Homework assignments (each Thursday)
  - **downloadable from StudIP after each lecture**
  - Tutorial groups will start **03.11**
- Homework is graded by our HiWis
  - **Send it by e-mail to your HiWi by the following week's Thursday no later than 15:00**
    - HiWi mail addresses are in StudIP
- Tutorial groups meet in room IZ 251



# 0. Homework and Tutorials

- **Groups of 3 students!**
  - Homework to be sent in **before the next lecture**
    - **3 pm sharp!** (~1 week of time)
  - Please send us **PDF files only!**
  - Mark each page with
    - **both your names** and
    - **matriculation numbers**
  - Adhere to our file naming guidelines
- Any **deviations** will lead to **0pts** for the sheet





# 0. Homework and Tutorials

- **So, who is my HiWi?**
  - We will open **Stud.IP groups** where you can assign yourself to a tutorial group, iff you do not yet have the Studienleistung
    - please make sure that **you and your partners are assigned to the same group** in Stud.IP!
    - groups are open starting **23.10 18:00 till 30.10 18:00**
  - The HiWi will:
    - send you the corrected homework before his tutorial starts
    - discuss the respective homework in his tutorial session



# 0. Tutorial Groups

Vorlesung: Relationale Datenbanksysteme 1 (WiSe 2025/26)

Verwaltung Übersicht **Teilnehmende** Dateien Ablaufplan



Lehrende

Nachname, Vorname

01 Balke, Wolf-Tilo

Mitwirkende

Teilnehmende

Gruppen

**Open: 23.10 18:00  
till 30.10 18:00**

2.

1.

3.

Vorlesung: Relationale Datenbanksysteme 1 (WiSe 2025/26)

Verwaltung Übersicht **Teilnehmende** Dateien Ablaufplan



Teilnehmende nach Gruppen

> Gruppe 01 - Montag 09:45 - 11:15 (0/18)

> Gruppe 02 - Montag 11:30 - 13:00 (0/18)



# 0. Classroom Exercise

- **Each Thursday** we will upload a classroom exercise consisting of:
  - **an exercise sheet**
  - **a solution**
  - **a recording to explain the solution**
- Practice for **yourself!** Try to solve the sheet and then watch the video. Additional service to practice for the exam and to solve the homework!



Institut für Informationssysteme  
Technische Universität Braunschweig

I. Große Übung


Data Modeling I

Hermann Kroll und Florian Plötzky  
Institute for Information Systems  
TU Braunschweig, Germany



# 0. Why should you be here?

- Its mandatory in your course of study....
- Database system are an **integral part** of most businesses, workflows and software products
- There is an abundance of **jobs** for people with good **database skills**
  - help yourself to put you into a good position within the job market
  - prepare for a sunny and wealthy future!





# 0. Why should you be here?

Job descriptions also exactly describe this course...

**bestica**

Database Analyst from Bestica  
Sierra Vista, AZ

Apply

Save Email

Details Highlights

Bestica is a trusted provider of solutions in Information Technology and Healthcare sectors to the DoD, Federal and Commercial markets. Our guiding principle and core values help us care for our people and the community; and build a culture of excellence which in turn helps us achieve our organizational objective of exceeding customer's expectations in each and every project. Headquartered in San Antonio, TX, we were founded in 2005.

**Resume check**

How do I match with this job?  
[Upload resume](#)

**Estimated Salary**

114,504 - 117,607 /year  
Database Analyst / Developer

**PCA**

Oracle Business Analyst I at Packaging C America  
Lake Forest, IL 60045

Details Highlights Company

As a Fortune 500 company, Packaging Corporation of America (PCA) is an ideas and sol business seeks to be the leader in helping our customers — large and small — package, transport and display products of all kinds. Our white paper business delivers Paper With Purpose by providing the highest level of customer service and operational excellence. We have approximately 15,000 team members in more than 100 locations in the United States that strive to meet the local needs of our customers. Our mission is to serve the needs of our customers, today and tomorrow, with products and services that exceed expectations for performance and environmental responsibility.

**Resume check**

How do I match with this job?  
[Upload resume](#)

**Estimated Salary**

116,799 - 130,221 /year  
Oracle Consultant / Analyst

monster  
<http://jobsearch.monster.com/>



# 0. Course Objectives

- After successfully completing this course you should be able to explain the fundamental terms of...
  - **databases** in general
  - the **relational model**
  - **theoretical and practical aspects of query languages**
  - conceptual and logical **design** of databases including **normalization**
  - application programming
  - further concepts like constraints, views, indexes, transactions and object databases



# 0. Course Objectives

- You should furthermore be able to
  - design and implement a database for any specified domain using **ER-Diagrams** or **UML-Diagrams**, the Relational Model and SQL-DDL
  - **normalize** a given relational database schema
  - **enhance** the database with views, indexes, constraints, and triggers
  - formulate data retrieval **queries** in **SQL**, Relational **Algebra**, and Relational **Calculi**
  - write programs **accessing databases** using JDBC




# 0. Contents of this Course

Lecture	Topic
1	Introduction
2	Data Modeling I
3	Data Modeling 2
4	View Integration
5	Relational Model
6	Relational Algebra
7	Relational Calculus
8	SQL I
9	SQL 2
10	Normalization
11	Application Programming I
12	Application Programming 2
13	Object Persistence
14	NoSQL



# 0. Recommended Literature



- **Fundamentals of Database Systems (EN)**
  - Elmasri and Navathe
  - Addison-Wesley
- **Database System Concepts (SKS)**
  - Silberschatz, Korth, and Sudarshan
  - McGraw Hill
- **Database Systems (GUW)**
  - Garcia-Molina, Ullman, and Widom
  - Prentice Hall
- **Datenbanksysteme (KE)**
  - Kemper, and Eickler
  - Oldenbourg





# 0. Recommended Literature

- **Database Modeling and Design: Logical Design**
  - Teorey, Lightstone, and Nadeau
  - Morgan Kaufmann
  - <http://www.sciencedirect.com/science/book/9780123820204>
- **SQL Cookbook**
  - Molinaro
  - O'Reilly
- **PostgreSQL Documentation**
  - <https://www.postgresql.org/docs/>
- **W3Schools SQL**
  - <http://www.w3schools.com/sql/>





# 0. Courses at ifis

- Basic course in databases
  - Relational Databases I (Bachelor)
    - What can we do with an DBMS?
    - Conceptual modeling, data retrieval, relational model, SQL, building applications, basic data models
  - Relational Databases II (Bachelor or Master)
    - How can we implement a DBMS?
    - Storage models, query optimization, transactions, concurrency control, recovery, data security





# 0. Courses at ifis

- Advanced courses in Information Systems (Master)
  - Information Retrieval and Web Search Engines
  - Multimedia Databases
  - Distributed Data Management
  - Knowledge-Based Systems and Deductive Databases
  - Data Warehousing and Data Mining Techniques





# I Introduction

- **What is a Database?**
- Characteristics of a Database
- History of Databases

All images, with the exception of some screenshots, book covers and logos are from unsplash.com, wikipedia.org or wikimedia.org unless stated otherwise.



PostgreSQL



MySQL®

ORACLE®

Microsoft®  
SQL Server®

SYBASE®



# I.I What is a Database?

- **Managing** large amounts of **data** is an integral part of most nowadays business and governmental activities
  - collecting taxes
  - bank account management
  - bookkeeping
  - airline reservations
  - human resource management
  - ...





# I.I What is a Database?

- **Databases** are needed to manage that **vast amount of data**
- A database (**DB**) is a collection of **related data**
  - represents some aspects of the **real world**
    - **universe of discourse**
  - data is logically **coherent**
  - is provided for an intended group of **users** and **applications**





# I.I What is a Database?

- As for today, the **database industry** is one of the most successful branches of computer science
  - constantly growing since the 1960s
  - more than **\$8 billion revenue** per year
  - DB systems found in nearly any application
  - ranging from large commercial transaction-processing systems to small open-source systems for your Web site





# I.I What is a Database?


- Databases are maintained by using a collection of programs called a database management system (**DBMS**), that deals with
  - definition of data and structure
  - physical construction
  - manipulation
  - sharing/protecting
  - persistence/recovery





# I.I File Systems

- A file system is not a database!
- File management systems are **physical** interfaces





# I.I File Systems

- **Advantages**
  - fast and easy access
- **Disadvantages**
  - uncontrolled redundancy
  - manual maintenance of consistency
  - limited data sharing and access rights
  - poor enforcement of standards
  - excessive data and access paths maintenance





# I.I Databases

- Databases are **logical** interfaces
  - retrieval of data using **data semantics**
  - controlled redundancy
  - data consistency & integrity constraints
  - effective and secure data sharing
  - backup and recovery
- However...
  - more complex
  - more expensive data access





# I.I Databases


- **DBMS** replaced previously dominant file-based systems in **banking** due to special requirements
  - **simultaneous** and quick access is necessary
  - failures and loss of data **cannot** be tolerated
  - data always has to remain in a **consistent** state
  - frequent queries and modifications





# I Introduction


- What is a Database?
- **Characteristics of a Database**
- History of Databases





## I.2 Characteristics of DBs


- Databases **control redundancy**
  - same data used by different applications or tasks is **stored only once**
  - access via a **single interface** provided by DBMS
  - redundancy only purposefully used to speed up data access (e.g. materialized views)
- Problems of **uncontrolled redundancy**
  - updating data may result in inconsistent data





# I.2 Characteristics of DBs

- Databases are **well-structured** (e.g. ER model)
  - simple banking system



- Relational Databases provide
  - **catalog** (data dictionary) contains all **meta data**
  - defines the **structure** of the data in the database



# I.2 Characteristics of DBs

- Databases support **declarative querying**
  - just specify what you want, not how and from where to get it
  - queries are separated and abstracted from the actual physical organization and storage of data
- Get the firstname of all customers with lastname *Smith*
  - file system: trouble with physical organization of data
    - Load file `c:\datasets\customerData.csv`.
    - Build a regular expression and iterate over lines:  
If 2<sup>nd</sup> word in line equals *Smith*, then return 1<sup>st</sup> word.
    - Stop when end-of-file marker is reached.
  - database system: simply query
    - `SELECT firstname FROM data WHERE lastname='Smith'`



# I.2 Characteristics of DBs

- Databases aim at **efficient** manipulation of data
  - physical tuning allows for good data allocation
  - indexes speed up search and access
  - query plans are optimized to improve performance
- Example: Simple Index

**Index File**  
(checking accounts)

number
4543032
7809849
8942214

Data File		
number	type	balance
1278945	saving	€ 312.10
2437954	saving	€ 1324.82
4543032	checking	€ -43.03
5539783	saving	€ 12.54
7809849	checking	€ 7643.89
8942214	checking	€ -345.17
9134354	saving	€ 2.22
9543252	saving	€ 524.89




# I.2 Characteristics of DBs

- **Isolation** between applications and data
  - database employs **data abstraction** by providing **data models**
  - applications work only on the **conceptual representation** of data
    - Data is strictly **typed** (Integer, Float, Timestamp, Varchar, ...)
    - Details on where data is actually stored and how it is accessed are **hidden** by the DBMS
    - Applications can access and manipulate data by invoking **abstract operations** (e.g. SQL statements)
  - DBMS-controlled parts of the file system are **protected** against external manipulations (tablespaces)



## I.2 Characteristics of DBs


- **Example:** Schema can be changed and tablespace moved without adjusting the Application's SELECT





# I.2 Characteristics of DBs

- **Example:** Schema can be changed and tablespace moved without adjusting the Application's SELECT





# I.2 Characteristics of DBs



- Supports multiple **views** of the data
  - views provide a different perspective of the DB
    - a user's conceptual understanding or task-based excerpt of the data (e.g. aggregations)
    - security considerations and access control (e.g. projections)
  - for applications, a view does not differ from a table
  - views may contain **subsets** of a DB and/or contain **virtual data**
    - virtual data is **derived** from the DB (mostly by simple SQL statements, e.g. joins over several tables)
    - can either be computed at query time or **materialized upfront**



# I.2 Characteristics of DBs

- Example views: **Projection**
  - *saving account clerk vs. checking account clerk*

**Original Table**

number	type	balance
1278945	saving	€ 312.10
2437954	saving	€ 1324.82
4543032	checking	€ -43.03
5539783	saving	€ 12.54
7809849	checking	€ 7643.89
8942214	checking	€ -345.17
9134354	saving	€ 2.22
9543252	saving	€ 524.89

**Saving View**

number	balance
1278945	€ 312.10
2437954	€ 1324.82
5539783	€ 12.54
9134354	€ 2.22
9543252	€ 524.89

**Checking View**

number	balance
4543032	€ -43.03
7809849	€ 7643.89
8942214	€ -345.17



# I.2 Characteristics of DBs

- **Sharing** of data and support for **atomic multi-user transactions**
  - transactions are a **series of database operations** executed as **one logical operation**
  - **concurrency control** is necessary for maintaining consistency
    - multiple users and applications may access the DB at the same time
  - **transactions** need to be **atomic** and **isolated** from each other





# I.2 Characteristics of DBs

- **Example:** Atomic transactions

- **Program:**

Transfer x Euros from Account 1 to Account 2

1. Debit amount x from Account 1
2. Credit amount x to Account 2





## I.2 Characteristics of DBs

- **Example:** Atomic transactions

- **Program:**

- Transfer x Euros from Account 1 to Account 2

- 1. Debit amount x from Account 1
    2. Credit amount x to Account 2

- **But what happens if the system fails after performing the first step?**





# I.2 Characteristics of DBs

- **Example:** multi-user transactions

- **Program:**

- Withdraw x Euros from Account I

1. Read *old balance* from DB
2. Set *new balance* to *old balance* – x
3. Write *new balance* back to the DB

- **Problem:** Dirty Read

- Account I has €500
    - User 1 deduces €20
    - User 2 deduces €80 **at the same time**


- Without multi-user transactions, Account I will have either €480 or €420, but not the correct €400
    - Both users read *old value* of €500 simultaneously, both deduce either €20 or €80, both write back new value (in random order)





# I.2 Characteristics of DBs

- **Persistence of data and disaster recovery**
  - data needs to be persistent and accessible at all times
  - **quick recovery from system crashes without data loss**
  - recovery from natural disasters (fire, earthquake, ...)





# I.2 Database Users

- Usually **several groups of persons** are involved in the daily usage of a large DBMS (many job opportunities for smart DB people...)
- Persons directly involved on DB level
  - **Database Administrators**
    - responsible for tuning and maintaining the DBMS
    - management of storage space, security, hardware, software, etc.
  - **Database Designers**
    - identify the data that needs to be stored and chooses appropriate data structures and representations
    - integrate the needs of all users into the design





# I.2 Database Users

## – Application Developers

- identify the requirements of the end-users
- develop the software that is used by (naïve) end-users to interact with the DB
- cooperate closely with DB designers



## – Data Analyst

- Analyzes trends in data to uncover valuable feedback for business operations
- i.e., discover sales trends, upcoming supply shortages, etc.

## • Persons working behind the scenes

### – DBMS Designers and Implementers

- implement the DBMS software

### – Tool Developers

- develop generic tools that extend the DBMS' functionalities

### – Operators and maintenance personnel

- responsible for actually running and maintaining the DBMS hardware





## I.2 Database Users

- Persons using the database (End Users)
  - All people who use the DB to do their job
- End Users can be split into
  - **Naïve End Users**
    - make up most DB users
    - usually **repeat** similar tasks over and over
    - are supported by predesigned interfaces for their tasks
    - e.g. bank tellers, reservation clerks, ...





# I.2 Database Users

## – Sophisticated End Users

- require **complex** non-standard operations and views from the DB
- are familiar with the facilities of the DBMS
- can solve their problems themselves, but require complex tools
- e.g. engineers, scientists, business analysts, ...



## – Casual End Users

- use DB only from time to time, but need to perform different tasks
- are familiar with query languages
- e.g. people in middle or senior management





# I Introduction

- What is a Database?
- Characteristics of a Database
- **History of Databases**





# I.3 History of DBs

*Detour*

- Databases have an exceptional history of development
  - many synergies between **academic**, **governmental** and **industrial** research
  - much to be learned from it
  - most popular concepts used today have been invented decades ago







# I.3 History of DBs

*Detour*

- The beginnings
  - 1880: U.S. Bureau of Census instructs **Herman Hollerith** to develop a machine for storing census data
  - result: **Punch card** tabulating machine
    - the evaluation of 1880's census took 8 years
    - 1890's has been finished after only 6 years
  - leads to the foundation of **IBM**
    - International Business Machines
  - data processing machines soon established in accounting





# I.3 History of DBs

*Detour*

- One of Hollerith's punch cards:

Lp	A	B	C	A	B	C	Lp	CH	%	Gn	Ag	Cx	Ct	SM	Ir	HM	WT	A	G	E	F	a	d
Ch	D	E	F	D	E	F	Lu	Ch	S	Se	Ma	Lb	FV	Or	Ca	X	Tb	B	D	X	*	b	*
Le	G	H	I	G	H	I	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
Cin	K	L	M	K	L	M	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
CS	N	O	P	N	O	P	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
LS	Q	R	S	Q	R	S	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
Kn	*	*	*	*	*	*	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4
RN	*	*	*	*	*	*	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5
QC	6	7	8	6	7	8	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
AV	*	1	m	*	1	m	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7
W	*	*	*	*	*	*	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
Sc	*	*	*	*	*	*	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9

3994



# I.3 History of DBs

*Detour*

- Tabulating machines
  - operations or “programs” directed by a plug board
  - up to 150 cards per minute
  - results were printed or punched for input to other processing steps





# I.3 History of DBs

*Detour*

- In 1951 IBM develops the electric **UNIVAC I**
  - first commercial computer produced in the USA
    - programmable (turing complete)
    - input (programs and data) with tape or punched cards
- In 1959, USA dominated the (still highly active) punch card machine market
  - within the USA, the Pentagon alone used more than 200 data processing computers, costing \$70 million per year





# I.3 History of DBs

*Detour*

- In **1964**, the term *data base* appeared for the first time in military computing using **time sharing systems**
  - data could be shared among users
  - but data was still bound to one specific application
    - similar data needed by multiple applications had to be duplicated
    - consistency problems when updating data
  - data structure highly-dependent on the hardware and (low-level) programming language used
    - inspired by punch cards and optimized for magnetic tapes
    - usually, no **relationships** between different records have been stored, just plain data





# I.3 History of DBs

*Detour*

- To turn stored data into a proper **database**, the following goals had to be achieved (McGee, 1981):
  - **Data Consolidation**
    - data must be stored in a central place, accessible to all applications
    - knowledge about relationships between records must be represented
  - **Data Independence**
    - data must be independent of the specific quirks of the particular low level programming language used
    - provide high-level interfaces to physical data storage
  - **Data Protection**
    - data must be protected against loss and abuse






# I.3 History of DBs

*Detour*

- **Data Consolidation** motivated the development of data models
  - Hierarchical Data Model
  - Network Data Model
  - **Relational Data Model**
  - Object-oriented Data Model
  - Semantic Data Model
- **Data Independence** inspired the development of query models and high-level languages
  - **Relational Algebra, SQL**
- **Data Protection** led to development of transactions, backup schemes, and security protocols






# I.3 History of DBs

*Detour*

- **Hierarchical Data Model**
  - first appearance in **IBM's IMS** database system, designed for the Apollo Program in **1966**
    - still, as of 2006, 95% of all Fortune 1000 companies used IBM IMS in their data backbone...
  - benefits from **advances in hardware design**
    - random access main memory and tape media available






# I.3 History of DBs

*Detour*

- **Hierarchical data model**
  - each type of record has some defined structured data
  - hierarchical **one-to-many** relationships
- **Advantages**
  - 1:n relationships can be expressed
  - can easily be stored on tape media
- **Disadvantages**
  - no n:m relationships
  - no Data Independence





# I.3 History of DBs

*Detour*

- **Network Data Model**

- in the mid-1960th, direct access storage devices (DASD) gained momentum
  - primarily hard disks
  - more complex storage schemes possible
- Hierarchical Data Model failed,  
e.g. for bill-of-material-processing (BOMP)
  - many-to-many relationships needed
  - development of the IBM DBOMP system (1960)
- result: Network Data Model
  - two types of files: master files, chain files
  - chain file entries could chain master file entry to one another



# I.3 History of DBs

*Detour*

- **Network Data Model**
  - the model was standardized by Charles W. Bachman for the **CODASYL** Consortium in 1969
    - CODASYL = Conference of Data Systems Languages
    - thus, also called the CODASYL model
  - allowed for more natural modeling of **associations**
- **Advantages**
  - **many-to-many-relationships**
- **Disadvantages**
  - no declarative queries
  - queries must state the data access path





# I.3 History of DBs

*Detour*

- **The relational data model**
- Published by **Edgar F. “Ted” Codd** in 1970, after several years of work
  - *A Relational Model of Data for Large Shared Data Banks*, Communications of the ACM, 1970
  - employee of IBM Research
    - IBM **ignored** his idea for a long time as not being “practical” while pushing its hierarchical IMS database system
    - other researchers in the field also **rejected** his theories
    - finally, he received the Turing Award in 1981






# I.3 History of DBs

*Detour*

- **Idea underlying the relational model:**
  - database is seen as a collection of **predicates** over a finite set of **predicate variables**
    - example
      - $\text{dislikes}(x, y)$
      - $\text{dislikes}(\text{'Ted Codd'}, \text{'hierarchical IMS database system'})$  (TRUE)
      - $\text{dislikes}(\text{'IBM'}, \text{'hierarchical IMS database system'})$  (FALSE)
    - the set of all true assignments is called a **relation**
    - relations are stored in **tables**
  - contents of the DB are a **collection of relations**
  - queries are also **predicates**
    - queries and data are very similar
    - Allows for **declarative querying**






# I.3 History of DBs

*Detour*

- It's really like a collection of index cards
  - more details during the next weeks...





# I.3 History of DBs

*Detour*

- Beginning 1977, **Lawrence J. Ellison** picked up the idea and created **Oracle DB**



- and became insanely rich – long time in the Top 10 of the richest people
- in 2015 Oracle ranked second on the list of largest software companies in the world, right after Microsoft





## I.3 History of DBs

*Detour*

- During the 1970s, IBM had also decided to develop a relational database system
  - **System R** with the first implementation of the **SQL** declarative query language (**SEQUEL**)
  - at first, mostly a research prototype, later became the base for **IBM DB2**
- Ingres Database was developed at UC Berkeley as a research project
  - It became PostgreSQL in 1996





# I.3 History of DBs

*Detour*

- Lately, the so-called NoSQL databases have become popular
  - NoSQL systems are mostly non-relational database systems
    - They might support SQL as a query language
  - Its development was driven by IT companies like Amazon, Google and Facebook/Meta
    - Databases have to deal with huge amounts of data and massive numbers of users
  - NoSQL comprise: key-value stores, document stores, column stores, distributed databases, graph databases, triple stores
    - For details, visit our advanced lectures



# I Summary

- Databases
  - are **logical interfaces**
  - support **declarative querying**
  - are **well-structured**
  - aim at **efficient manipulation** of data
  - support **control redundancy**
  - support **multiple views** of the data
  - support **atomic multi-user transactions**
  - support **persistence** and **recovery** of data





# 2 Coming soon...

- **Next Lecture**
  - Phases of DB Design
  - Data Models
  - Basic ER Modeling
    - Chen Notation
    - Mathematical Model

