



ifis

Institut für Informationssysteme
Technische Universität Braunschweig

Relational Database Systems I

Wolf-Tilo Balke


Niklas Kiehne, Enrique Pinto Dominguez

Institut für Informationssysteme
Technische Universität Braunschweig
<http://www.ifis.cs.tu-bs.de>



4 View Integration

- **View integration**
- Resolving conceptual incompatibility
- Entity clustering for ER models
- Commercial dimension:
The BEA story





4. I Business Integration


- Business currently is a world of M&A
 - companies need to diversify/enhance their portfolio
 - but it is expensive to develop necessary applications
 - knowledge gathering costs time
 - will the output be worth it?
 - **idea:** rely on people who are **already knowledgeable**
 - acquire small, specialized, and promising companies
 - merge with big players in the field





4. I Business Integration

- Examples
 - the Daimler-Chrysler merger
 - the Oracle-Sun merger
 - Oracle buys PeopleSoft,
Siebel Systems, BEA Systems, ...
 - Siebel Sales as CRM tool now part of
Oracle's business intelligence suite



DAIMLERCHRYSLER

ORACLE
 Sun
microsystems

PeopleSoft

SIEBEL.

bea®
Think liquid.™



4. I Business Integration

- Merged (parts of) businesses are administrated by
 - different specialized software systems?
 - one company-wide system?
- Usually, there is a **historical evolution** of separate tools and programs
 - e.g., accounting, sales & marketing, development
 - based on individual requirements
- However, often a **unified view** is needed
 - e.g., for business intelligence and warehousing
 - Data Warehousing is also a **great lecture** at ifis...



4. I View Integration

- Usually, there are **several conceptual schemas**
 - **several designers** are part of the modeling process (modular software development)
 - **different tasks** were modeled within the same organization (legacy systems)
 - **several organizations** need to be integrated (business integration)





4. I View Integration

- **View Integration**
 - several conceptual schemas need to be combined into a **unified global schema**
 - all differences in perspective and terminology have to be resolved
 - all redundancy has to be removed
- **But,... what happens,
if you don't integrate?!**






4. I The Mars Desaster

Detour

- **Example:** Big NASA project **Mars Surveyor**
 - the 1998 missions investigated *Volatile and Climate History on Mars*
 - characterization of climate change and its evolving impact on the distribution of water
 - **idea:** explore the **polar ice caps** of Mars and see whether there is water ice
 - about three months of experiments on Mars were scheduled






4. I The Mars Desaster

Detour

- Two vehicles:
Mars climate orbiter and **Mars polar lander**
 - the lander was supposed to probe the layers of ice and dust on the polar ice caps to investigate changes
 - the orbiter was built to monitor the daily weather and record changes in water vapor and dust in the atmosphere





4. I The Mars Desaster

Detour

- Catastrophic failure
 - the Mars climate orbiter approached Mars up to 57 km instead of 150 km, and was **destroyed** in the atmosphere on September 23, 1999
 - the Mars polar lander **crashed** during its attempted landing on Mars, December 3, 1999
 - **\$327.6 million** in total for both
 - \$193.1 million for spacecraft development
 - \$ 91.7 million for launch
 - \$ 42.8 million for mission operations






4. I The Mars Desaster

Detour

- Why did the **climate orbiter** come too close to Mars' atmosphere?
 - **many organizations** were involved in the development
 - there was no global schema
 - navigation software produced by Lockheed Martin used **non-metric** units (i.e. inches, feet, and pounds)
 - NASA used **metric** units
 - a small correction of the course led to the fatally low orbit...






4. I The Mars Desaster

Detour


- Happy ending!
 - the next try was the successful **2001 Mars Odyssey**
 - the measurements pointed to **water ice on Mars**
 - confirmed by the Mars Express (ESA) in 2004
 - the polar caps consist of 85% carbon dioxide (CO_2) ice and 15% water ice





4 View Integration

- View integration
- **Resolving conceptual incompatibility**
- Entity clustering for ER models
- Commercial dimension:
The BEA story





4.2 Resolving Incompatibilities


- **Schema diversity** occurs when different users develop their own understanding of the world
 - the same reality is not always modeled in the same way due to different information needs or workflows
- **Common principle** for schema integration
 - identify the parts of the input schemas that represent the same reality
 - unify their representations






4.2 Basic Steps

- There are **four basic steps** needed for conceptual schema integration
 1. pre-integration analysis
 2. comparison of schemas
 3. conformation of schemas
 4. merging and restructuring of schemas
- The integration process needs continual refinement and reevaluation





4.2 Schematic View



pre-integration analysis

comparison of schemas


conformation of schemas

merging and restructuring



4.2 Pre-Integration Analysis


- **Pre-integration analysis** takes a close look on the individual conceptual schemas to decide for an **adequate integration strategy**
 - the larger the number of constructs, the more important is modularization
 - is it really sensible/possible to integrate all schemas?





4.2 Pre-Integration Analysis

- First, an **integration strategy** has to be chosen
- Schema integration can either be performed **many at a time**, ...
 - requires **only one consistent merge**
 - conflict analysis from many schemas is difficult
 - may **conserve efforts** if done well





4.2 Pre-Integration Analysis

- ... or can be performed **two at a time**


- results are step by step accumulated into a single schema

- how to choose **the right order** of schema comparisons

- the order can influence the final result

- selecting the first schema

- mixed strategy: skeleton schema
 - otherwise: most important schema





4.2 Comparison of Schemas

- The **resolution of conflicts** needs a thorough comparison of schemas
 - general question: how do entities **correspond?**
 - **naming conflicts** can be detected, e.g. by scanning the data dictionary
 - **structural conflicts** regarding semantics have to be resolved
 - different cardinalities in relationships
 - key conflicts
 - ...






4.2 Comparison of Schemas

- The individual **perspective of the world** and the **level of abstraction** are major reasons for conceptual incompatibilities

- Example

- A **product** is a **unit of sale** for the marketing department, but consists of **parts** in the view of the engineering department.





4.2 Comparison of Schemas


- The **level of abstraction** directly influences the schema design
- **Simple example:** A customer buys a product
- The **marketing view** focuses on how many people buy some product, e.g., for advertising
 - only the characteristics of the customer and product and the connection are needed





4.2 Comparison of Schemas


- The **accounting view** also needs the exact order number for identification of individual customer transactions
 - focus is on the purchase, but individual orders have to be distinguished





4.2 Comparison of Schemas

- The **sales view** needs all individual order details, e.g., for troubleshooting or CRM
 - focus is on orders
(which provide the basis for purchase contracts)





4.2 Comparison of Schemas

- Marketing



- Accounting




- Sales





4.2 Comparison of Schemas


- Different user groups use different names to refer to the same entities (differing terminology)
 - **synonyms:** two terms for the same entity
 - **homonym:** the same term for different entities
- **Rule of thumb:**
eliminate synonyms,
rename homonyms!





4.2 Conformation of Schemas

- The main goal is to make schemas **compatible** for integration
- Conformation usually needs **manual** interaction
 - conflicts need to be resolved semantically
 - rename entities/attributes
 - convert differing types, e.g.
convert an entity to an attribute
or a relationship
 - align cardinalities/functionalities
 - align different data types






4.2 Conformation of Schemas

- Besides renaming and type conversions, **abstraction** can be useful
 - generalization and aggregation allows to create new supertypes or subtypes
- Also **assertions and constraints** must be generalized or distributed among the type hierarchy
 - for example, checking accounts and saving accounts are both types of accounts, but may differ with respect to the minimum balance constraint



4.2 Conformation of Schemas

- **Example:** Resolving differing terminology
- **Homonyms:** Schema 1 and 2 both contain the term *jaguar*, but mean different entities
 - Rename to *jaguar_car* and *jaguar_animal*
- **Synonyms:** Schema 1 contains the term *jaguar*, whereas schema 2 contains the term *panther*
 - global schema should model *panther* *is_a* *jaguar*
 - constraint on the black color should be added





4.2 Merging and Restructuring

- How to merge schemas into a **global schema**?
 - copy all distinct **entities** from the individual schemas
 - apply renaming, overlapping entity integration, abstraction, attribute type conversions, etc.
 - put in the distinct **relationships** from all schemas
 - again use renaming, cardinality/functionality conversions, etc.
 - restructure the resulting **global schema**







4.2 Merging and Restructuring

- The final **restructuring of the schema** is driven by the goal of completeness, minimality, and ease of understanding
 - **completeness** mandates that all concepts in the global schema appear *semantically intact*
 - all different concepts of every individual schema are also part of the global schema
 - for each concept, there are no missing attributes, no constraints that cannot be met by all members of a type, etc.



4.2 Merging and Restructuring

- **minimality** enforces to remove all redundant concepts from the global schema
 - e.g. overlapping entities or redundant relationships
 - often, the question of minimality can only be decided **semantically**





4.2 Merging and Restructuring

– **ease of understanding** means that the global schema makes sense to the users

- in particular, abstraction and fine granular levels for entities can be very confusing
- Example: *Subtype entities have to be clearly distinguishable, and should have only attributes that are not inherited from the supertype.*





4.2 Example: View Integration *Detour*


- **Doomsday Legion (DDL)**
 - cooperation of villains from all over the world striving for global domination
 - channeling resources, staff, experience and power for reaching their goals
 - centralized and coordinated management of all shared assets
 - lairs
 - minions
 - assault squads
 - ...





4.2 Example: View Integration *Detour*


- Doomsday Legion schema





4.2 Example: View Integration *Detour*


- **Justice League (JL)**
 - federation of super-powered heroes fighting against global crime and villainousness
 - in particular: opposing the Doomsday Legion
 - central management of joint operations and resources





4.2 Example: View Integration *Detour*

- Justice League schema







4.2 Example: View Integration *Detour*

- And then, strange and evil **aliens invade earth** without any obvious reason
 - Justice League wants to save earth (that's what they do)
 - Doomsday Legion wants to save earth (without people, global domination is no fun)
 - great idea: **Join Forces**
 - *Defenders of the Earth*
 - great problem: joining large organizations is **not that easy**
 - beside the problem of ignoring old hatred, the data **schemas need to be integrated** for central mission control and planning



4.2 Example: View Integration *Detour*

- How to integrate?





4.2 Example: View Integration *Detour*

- Integrating the person models




- different structure
 - DDL more general → Merge JL into DDL
 - generalize *Hero* and *Sidekick* into *Person*
 - but: **Attribute Homonyms!**
 - DDL uses the real name of *name*, the villain identity is *alias*
 - JL puts real name into *secret identity* and hero name into *name*
 - name *Victor von Doom* and alias *Dr. Doom* vs.
name *Invisible Woman* and secret identity *Susan Storm*
 - attributes need to be **renamed** and **transformed** correctly



4.2 Example: View Integration

Detour

- Integrating the person models




– semantic consolidation

- Hero* and *Villains* should be treated the same
 - both are highly skilled and powerful super members of DotE
 - merge classes into *SuperOperator* class
- Sidekicks* and *Henchmen* are close Assistants of an operator
 - Heroes usually only have **one** Sidekick
 - Use more general 1:N association to also capture *Henchmen*



4.2 Example: View Integration *Detour*

- Integrating the person models



- contains *Heroes* and *Villains*, as well as their respective *Sidekicks* or *Henchmen*
- Heroes* and *Sidekicks* get an additional id




4.2 Example: View Integration

Detour

- **Integrating bases**

- Villains only have 3 types of bases, explicitly modeled
- Heroes may have any kind of base, given by the type attribute
- Two solutions
 - merge DDL into JL
 - geoOrbit, city, and coordinates become location
 - type is given by subclass
 - only possible in a lossless fashion because subtypes don't have additional attributes
 - merge JL into DDL
 - depending on type, a base is assigned to one of the subclasses
 - 4th subclass necessary for all other types (could also be merged into superclass)


BaseOfOperations
+name
+location
+type





4.2 Example: View Integration *Detour*

- Integrating bases



– only *Villains* owned lairs; no information of ownership for former *Hero* bases



4.2 Example: View Integration *Detour*

- Integrating skills and powers

- JL only stores *Super Powers*
- DDL stores all *Skills* (including super powers)
 - more general
- Merge JL into DDL
 - all old justice league *Super Powers* become *Skills* of the type *super power*
 - *name* is either null or manually completed
 - No information on *Skills* of Sidekicks


SuperPower
+description

Skills
+type
+name
+description



4.2 Example: View Integration *Detour*

- Integrated schema





4.2 Outlook

- View integration is a **semantic process**
 - this usually means a lot of **manual work**
 - computers can support the process by **matching** some (parts of) schemas
- There have been some approaches towards **(semi-)automatic matching** of schemas
 - matching is a **complex process** and usually only focuses on simple constructs like
Are two entities semantically equivalent?
 - the result is still rather error-prone...



4.2 Outlook

- **Basic methods** (that can of course be mixed freely)
 - **label-based matching**
 - for each label in one schema, consider all labels of the other schema and every time gauge their semantic similarity
 - **instance-based matching**
 - looking at the instances (of entities or relationships) one can e.g. find correlations between attributes
 - Are there duplicate tuples? or*
 - Are the data distributions in their respective domains similar?*
 - **structure-based matching**
 - abstracting from the actual labels, only the structure of the schema is evaluated, e.g. regarding element types, depths in hierarchies, number and type of relationships




4.2 Outlook

- Sometimes schema integration is **query-driven**
 - the integration is only needed in order to query several different information sources having different schemas
- In that case only a **schema mapping** is needed
 - basically the mapping is a **list of correspondences** between equivalent entities or relationships of heterogeneous schemas
 - the query can then be **translated** for each different schema using the mapping
 - the mapping can be derived manually or automatically from a respective matching



4 View Integration

- View integration
- Resolving conceptual incompatibility
- **Entity clustering for ER models**
- Commercial dimension: ...
The BEA story






4.3 Entity Clustering

- Sample schema integration result

National
Contaminants
Information
System (NCIS)

- © Fisheries and Oceans, Canada






4.3 Entity Clustering

- When multiple schemas are merged, global schema can become **very large**
 - many different entities and relationships between them
 - e.g. *global view of a company with all its dependencies*
 - but some parts from different views are entirely independent
 - e.g. *accounting does not need technical specifications of products*
- **Idea:** Cluster **semantically coherent parts** and abstract from their actual entities in the global schema



4.3 Entity Clustering

- Abstracting complex units allow showing the entire model on a **single sheet of paper**
 - easy to get an overview and easier to integrate units separately
 - zoom in** for more details






4.3 Clustering Concepts

- **Grouping** is an operation that combines entities and their relationships to form higher-level constructs
 - groups are called **entity clusters**
 - can also be performed **hierarchically** from the entire database (root entity cluster) over several levels down to the individual entities
 - all original entities are on clustering level I



4.3 Clustering Concepts

- Usually, **entity clusters** are depicted similar to normal entities in ER diagrams
 - by a **dark-bordered box**
 - numbered according to the **clustering level** and an identifier
 - interfaces for **inter-cluster** relationships have to be annotated





4.3 Grouping Operations


- **Grouping operations** are the fundamental components of entity clustering
 - all operations are **heuristic** in nature
- **Often occurring** operations are
 - dominance grouping
 - abstraction grouping
 - constraint grouping
 - relationship grouping
- They can be applied **recursively** or in a variety of combinations to produce higher level clusters





4.3 Grouping Operations


- **Dominance grouping** focuses on semantically dominant entities in the ER diagrams
 - hubs for **otherwise unconnected** or **weak** entities





4.3 Grouping Operations


- **Abstraction grouping** clusters entities of a specific super-type
 - especially helpful, if subclasses have no individual relationships





4.3 Grouping Operations


- **Constraint grouping** clusters entities related by the same constraint
 - e.g. integrity constraints such as XOR constraints





4.3 Grouping Operations

- **Relationship grouping** focuses on ternary or higher-degree relationships
 - the relationship is represented as a **whole**





4.3 Clustering Technique

- Identify all major functional areas and subareas in a top down analysis
 - functional areas are often defined during the requirement analysis as important organizational units (e.g. HR or R&D) or business activities
 - usually there will be a certain degree of overlap, for example employee data will be administrated by HR, but may also be needed in other departments



4.3 Clustering Technique

- The actual clustering has **four steps**
 - define points of grouping within each functional area
 - locate **dominant entities**, consider **abstraction**, find **n-ary** or **constrained relationships**, etc.
 - if such points do not exist, consider grouping the entire area
 - form entity clusters
 - use the **basic grouping operations** on elementary entities and their relationships to form higher level clusters
 - since entities might belong to several clusters, **define priorities** like *always prefer abstraction grouping, avoid crossing boundaries of functional areas, or leave entities ungrouped, if they belong to two or more groups at the same level of precedence*



4.3 Clustering Technique

- form higher level entity clusters
 - apply the grouping operations **recursively** to any combination of elementary entities and entity clusters
 - stop, if the diagram's complexity is **sufficiently low**:
This defines the root entity cluster
- validate the cluster diagram
 - check for **consistency** of the interfaces (relationships) between entities or entity clusters at each level of the diagram
 - **verify the meaning** of each level with the intended users



4 View Integration

- View integration
- Resolving conceptual incompatibility
- Entity clustering for ER models
- **Commercial dimension:**
The BEA story





4.4 The BEA Story

Detour

- What happens, if you don't integrate properly?
 - think about the Mars disaster...
- What happens, if you integrate?
 - well, your processes are improved and you become more efficient...
- What happens, if you **help others** to integrate?
 - short version: you found a company, get insanely rich and are finally bought by Oracle for 8.5 billion USD in 2008





4.4 The BEA Story

Detour

- **BEA Systems Inc.**

- founded in 1995 in San José, CA, USA
- before Oracle's takeover, the company had more than 4000 employees and about **one billion** in revenues
- **product lines**
 - Tuxedo for **distributed transaction processing** (1995)
 - WebLogic provides a **J2EE enterprise infrastructure** (1998)
 - AquaLogic provides a **service-oriented infrastructure** (2005)
- **acquisitions** of some companies specializing in middleware and business process management
 - e.g. WebLogic (1998), SolarMetric (2005), Plumtree Software (2006), Fuego (2006)





4.4 The BEA Story

Detour

- What are they actually doing?
 - case study

„DekaBank“

 Finanzgruppe

- the **DekaBank Group** is the central asset manager of the Sparkasse Financial Group managing funds of around 90 billion EUR
- the Bank wanted an **access layer** to central data sources so that all data for the portfolio structure is available for fund management
- in 2006 DekaBank deployed the **BEA AquaLogic Data Services Platform**, which models the central data uniformly in a technical context and provides these business objects to local applications in real time



4.4 The BEA Story

Detour

- two **Challenges**
 - consolidation of various pieces of information from numerous channels
 - provide the information in different formats such that local applications can further process the data
- finally, after a lot of integration, data is presented to the outside via a **standard access layer** in real time
- **duration:** about five month
- **costs:** ???





4.4 The BEA Story

Detour

- **BEA AquaLogic Data Services**

- special Feature: **easy-to-use modeling**

- *In an SOA environment, a data model must be flexible so that it can represent any complex entity and rich enough to provide information about data structure, relationships, and services to read or update.*
 - *Data services are illustrated in model diagrams and can easily be shared with others in the enterprise for greater data consistency and reuse.*
 - *Mappings and transformations can be designed in an easy-to-use GUI tool using a library of over 200 functions. For complex mappings and transformations, architects and developers can bypass the GUI tool and use an XQuery source code editor to define or edit services.*





4.4 The BEA Story

Detour

- What tools are actually given to support integration?
 - Data Translation Tool
 - transforms binary data into XML
 - transforms XML to binary data
 - Data Transformation Tool
 - transforms an XML to another XML
 - idea
 - transform data to application specific XML
 - transform to other application's XML or general schema
 - transform back to binary
 - **note:** the integration work still has to be done **manually**





4.4 The BEA Story

Detour

- *I cannot afford expensive BEA consultants and the AquaLogic Integration Suite, what now?*
 - do it all by **yourself**
 - most used technologies can be found as open source projects (data mappers, XSL engines, XSL editors, etc.)
 - do it **yourself with specialized tools**
 - many companies and open source projects are specialized in developing data integration and transformation tools
 - CloverETL
 - Altova MapForce
 - BusinessObjects Data Integrator
 - etc...






4.4 The BEA Story

Detour

- **Altova MapForce**

- same idea as the BEA Integrator
 - also based on XSL and a data description language
- editors for binary/DB to XML mapping
- editor for XSL transformation
- automatic generation of data sources, web-services, and transformation modules in Java, C#, C++






4.4 The BEA Story

Detour

- **CloverETL**

- based on own ETL transformation language
- core tools are open source
 - server and GUI tools are sold under commercial license
- can read data from any database
- (visually designed) ETL Script converts data into other data
 - XML
 - DB with different schema
 - etc.





4 Next Week

- Basic set theory
- Relational data model
- Transformation from ER
- Integrity Constraints
- From Theory to Practice

