



**ifis**

Institut für Informationssysteme  
Technische Universität Braunschweig

# **Relational Database Systems I**

**Wolf-Tilo Balke**

**Niklas Kiehne, Enrique Pinto Dominguez**

Institut für Informationssysteme  
Technische Universität Braunschweig  
<http://www.ifis.cs.tu-bs.de>



# 9 SQL 2

- **SQL data definition language**
- SQL data manipulation language  
(apart from **SELECT**)
- $\text{SQL} \neq \text{SQL}$
- Some advanced SQL concepts





## 9.1 Recap

## Summary

- Last week, you learned how to query an existing relational database

**SELECT** <attribute, function, scalar subquery>  
**FROM** <table, table subquery>  
[**WHERE** <condition>]  
[**GROUP BY** <attribute list>]  
[**HAVING** <condition>]  
[**UNION/INTERSECT/EXCEPT** <query>]  
[**ORDER BY** <attribute list>]



# 9.1 SQL DDL



- What's missing?

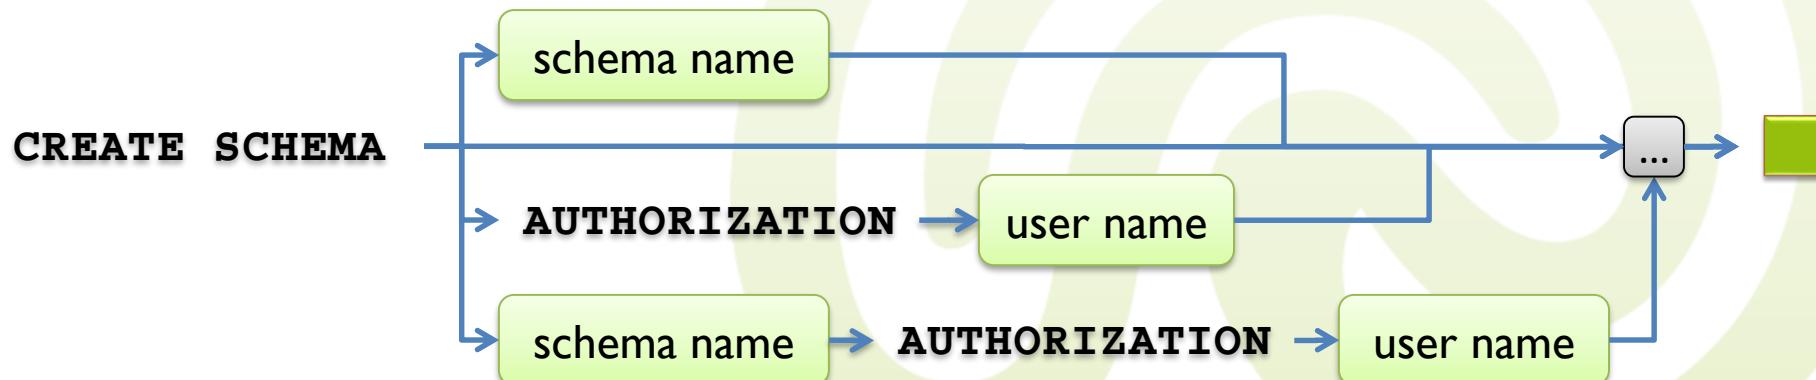
- how to **create** schemas, tables, ...
- how to **drop** schemas, tables, ...
- how to **alter** schemas, tables, ...
- how to **insert** new tuples into existing tables?
- how to **delete** tuples from existing tables?
- how to **update** tuples in existing tables?

} **DDL**  
} **DML**



# 9.1 SQL DDL: Schemas

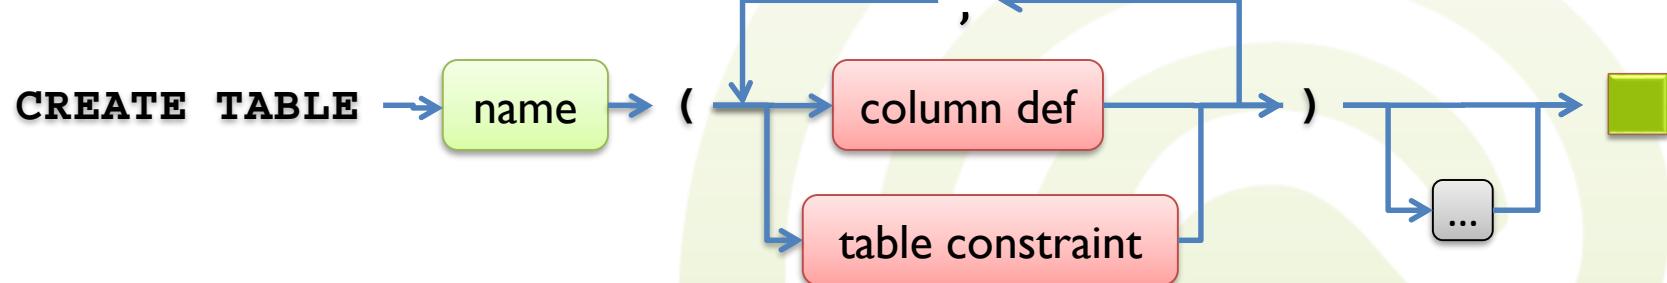
- **CREATE SCHEMA** creates a **new schema** with a given name for a given **owner**
  - if no schema name is provided, the **current username** is used
  - if no explicit owner is provided, also the **current user** is used
- Example
  - **CREATE SCHEMA allot\_club AUTHORIZATION karl**





# 9.1 SQL DDL: Tables

- **CREATE TABLE** creates a **new table** with a given name
  - contains column definition for **each column**
  - contains additional table-specific **structural constraints**





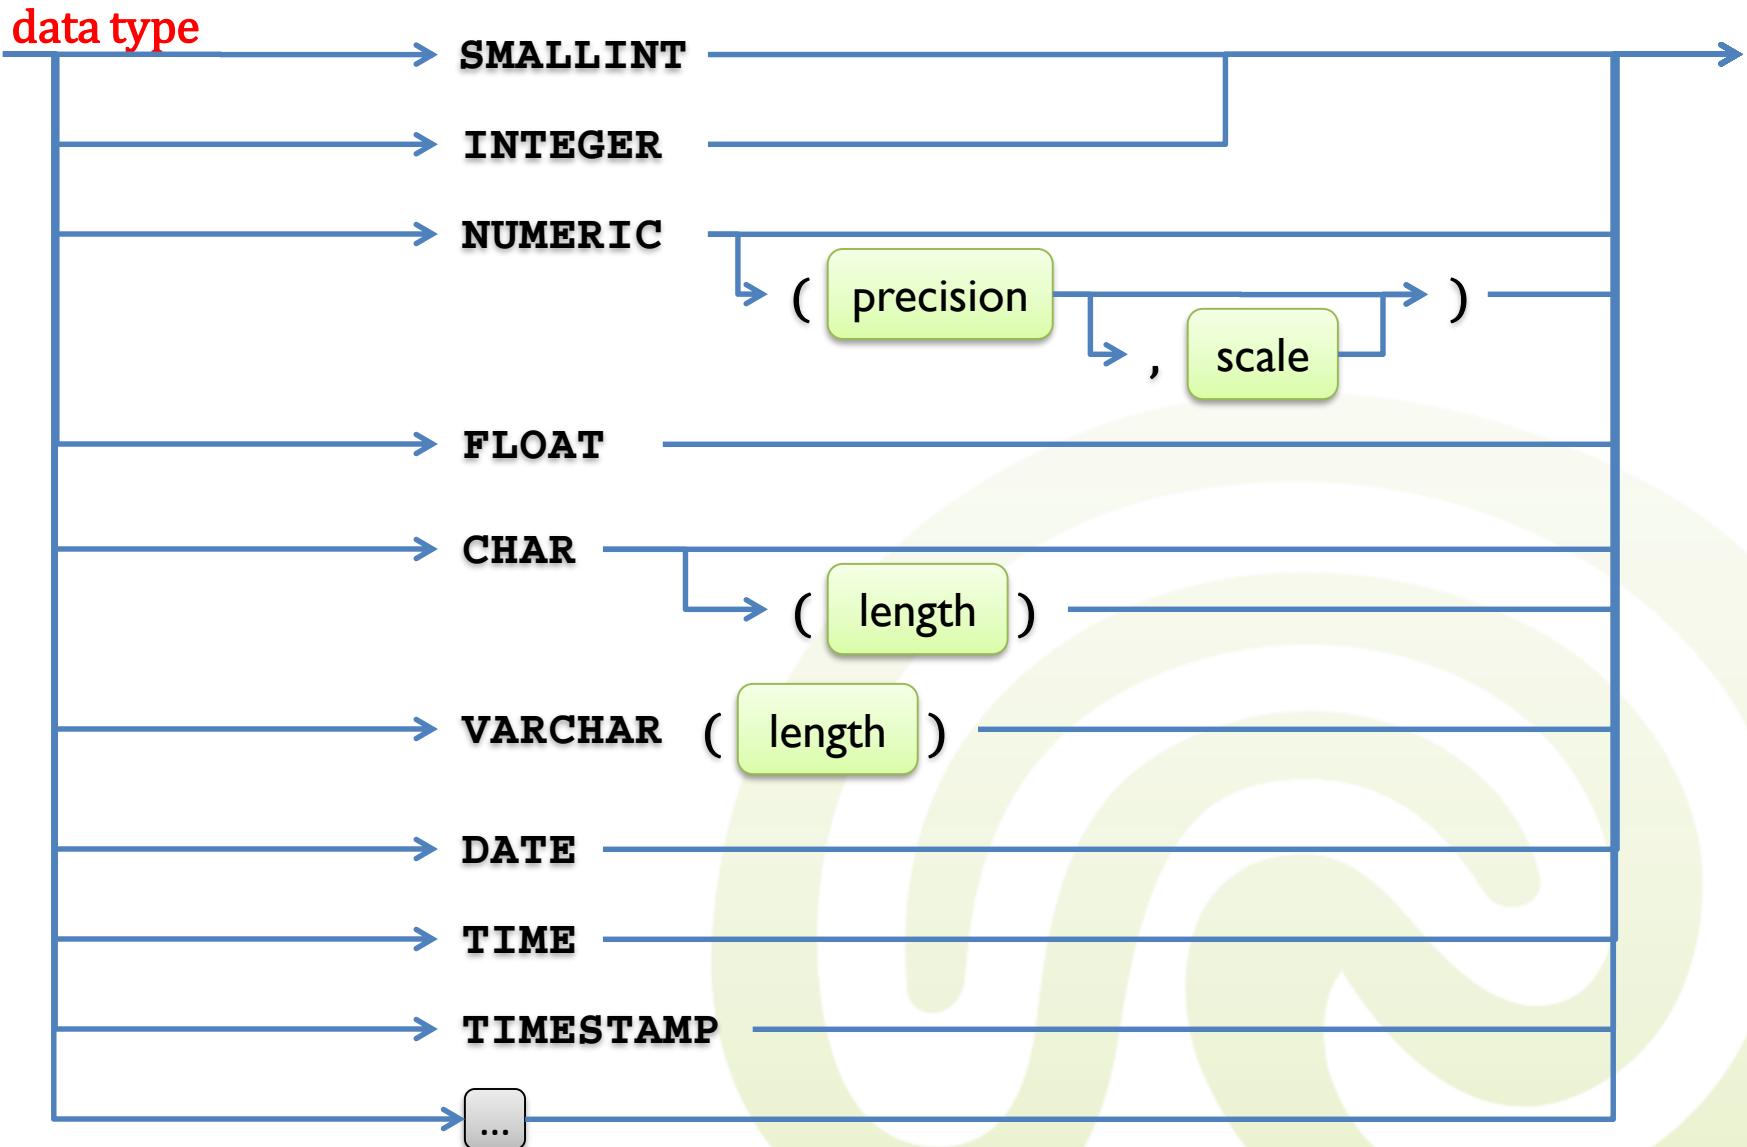
# 9.1 SQL DDL: Tables

- each column has a **name** and a **data type**
- each column may have multiple **column options**
- example
  - `CREATE TABLE member (  
 name VARCHAR(200),  
 age INTEGER  
)`





# 9.1 SQL DDL: Tables



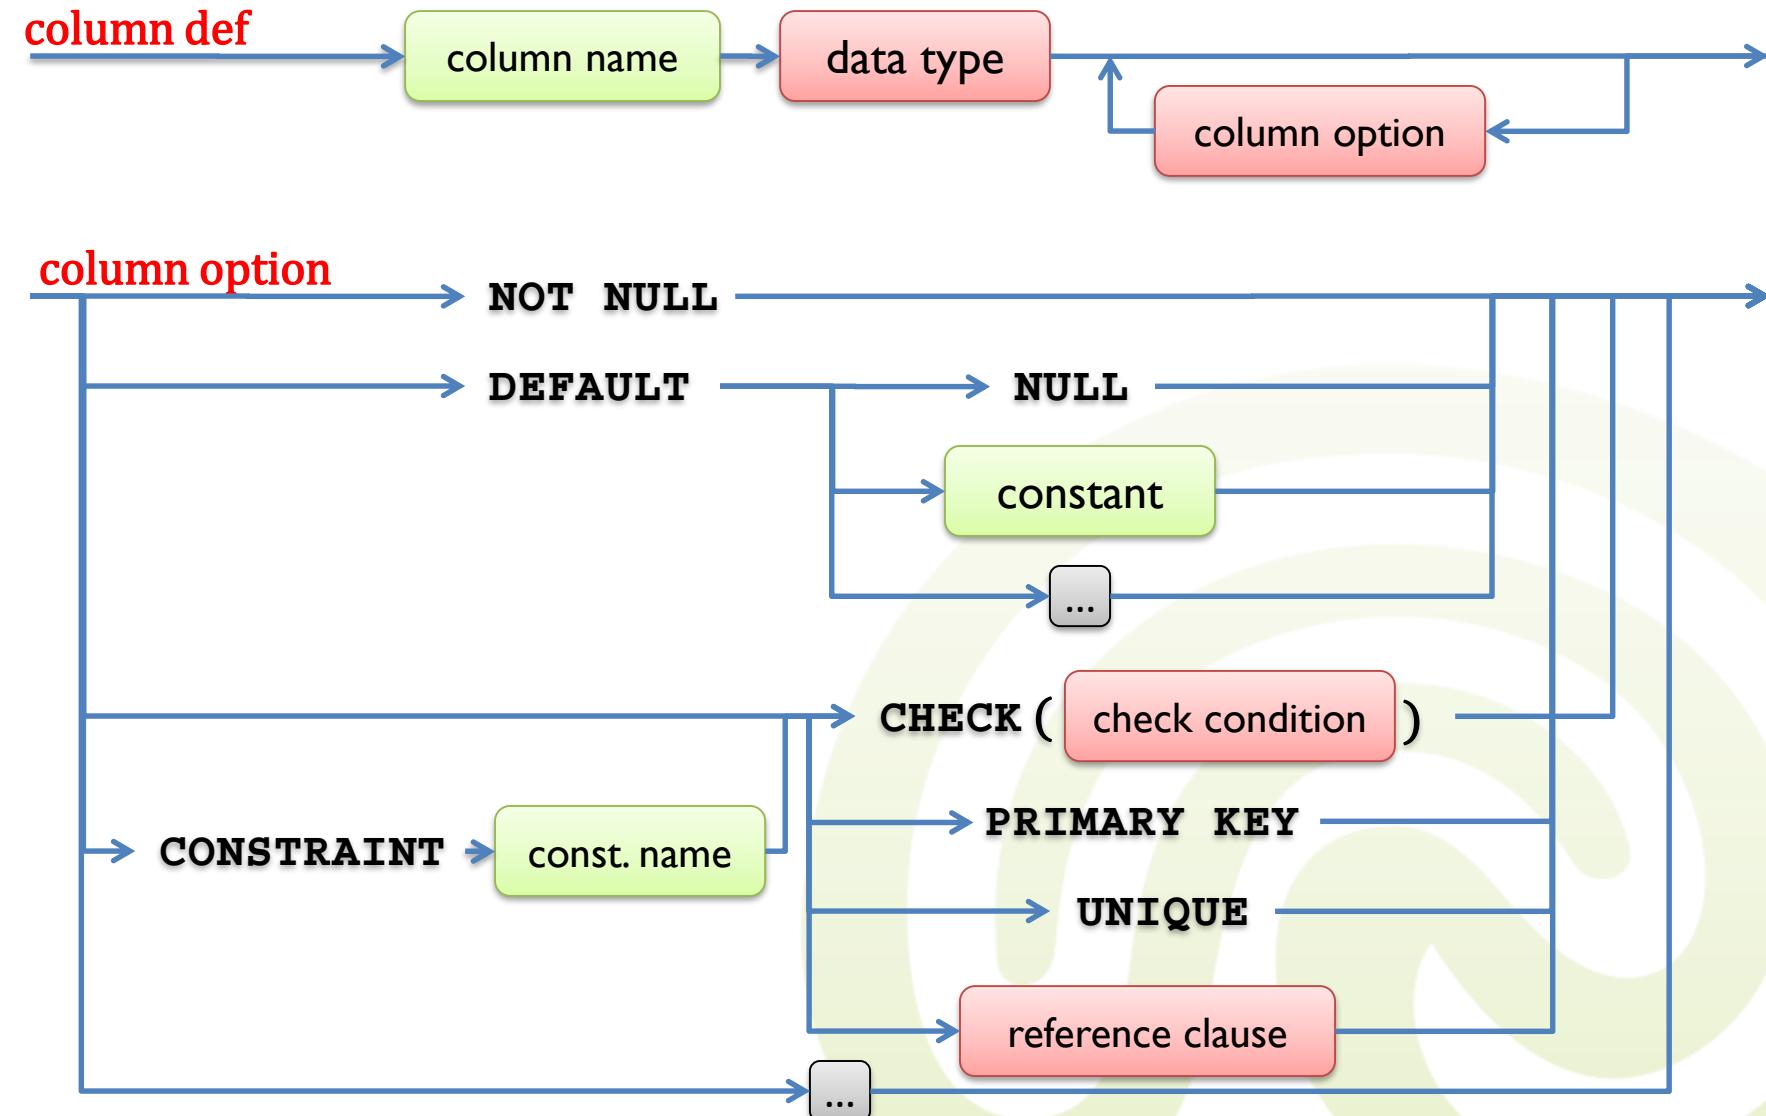


# 9.1 SQL DDL: Tables

Name	Syntax	description
Integer	INTEGER	Signed four-byte integer
Float	FLOAT/REAL/ DOUBLE PRECISION	Floating point number of approximate precision (the supported precision is implementation-dependent)
Numeric	NUMERIC (p , s)	An <b>exact</b> decimal number with $p$ digits: ( $p-s$ ) before the decimal point, and $s$ digits after the decimal point
Character	CHAR (x)	A textual string of fixed length $x$
Character varying	VARCHAR (x)	A textual string of length at most $x$
Date	DATE	Year, month, and day
Time	TIME	Hours, minutes, and seconds
Timestamp	TIMESTAMP	A date and a time



# 9.1 SQL DDL: Tables





# 9.1 SQL DDL: Tables

## – NOT NULL:

the **NULL** value is not allowed for the column

## – example

- **CREATE TABLE** member (  
    name **VARCHAR**(200) **NOT NULL**,  
    age **INTEGER** **NOT NULL**  
)



## – DEFAULT:

defines the default value if a value is not explicitly set

- usually a constant or **NULL**
- if omitted, **NULL** is the default

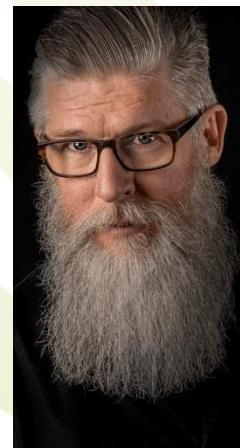


# 9.1 SQL DDL: Tables

- **Column constraints**

- restrict possible values for the current column
- may have a **unique name** indicated by  
**CONSTRAINT <name>**
  - if name is omitted, system creates a default name
- **CHECK:** user-defined constraint.  
To be valid, values have to satisfy the condition.
- example

- ```
CREATE TABLE person (
    name VARCHAR(200),
    age INTEGER CONSTRAINT adult
        CHECK (age >= 18)
)
```





# 9.1 SQL DDL: Tables

- **UNIQUE**: no duplicate values are allowed within this attribute
  - For multiple attributes, you need a different option (later)
  - NULL values are not considered equal
- example
  - **CREATE TABLE** person (  
    name **VARCHAR**(200) **NOT NULL UNIQUE**,  
    age **INTEGER NOT NULL**  
)





# 9.1 SQL DDL: Tables

- **PRIMARY KEY:** each table may have a **primary key** (optionally, but recommended) made up of at least one column
  - this option can only be used if the primary key consists of only one column
  - for multi-column primary keys, you need a different option (later)
  - implies **NOT NULL** and **UNIQUE**
- additionally, a **referential clause** may be specified (see next slides)





## 9.1 SQL DDL: Referential Integrity

- Rows in tables may **refer** to rows in other tables to capture **relationships**
- Of course, you should not be allowed to refer to a non-existing row
  - **referential integrity** between **primary keys** and **foreign keys** ensures that references are correct

| artist_id | artist_name      |
|-----------|------------------|
| 1         | Bono             |
| 2         | Cher             |
| 3         | Nuno Bettencourt |

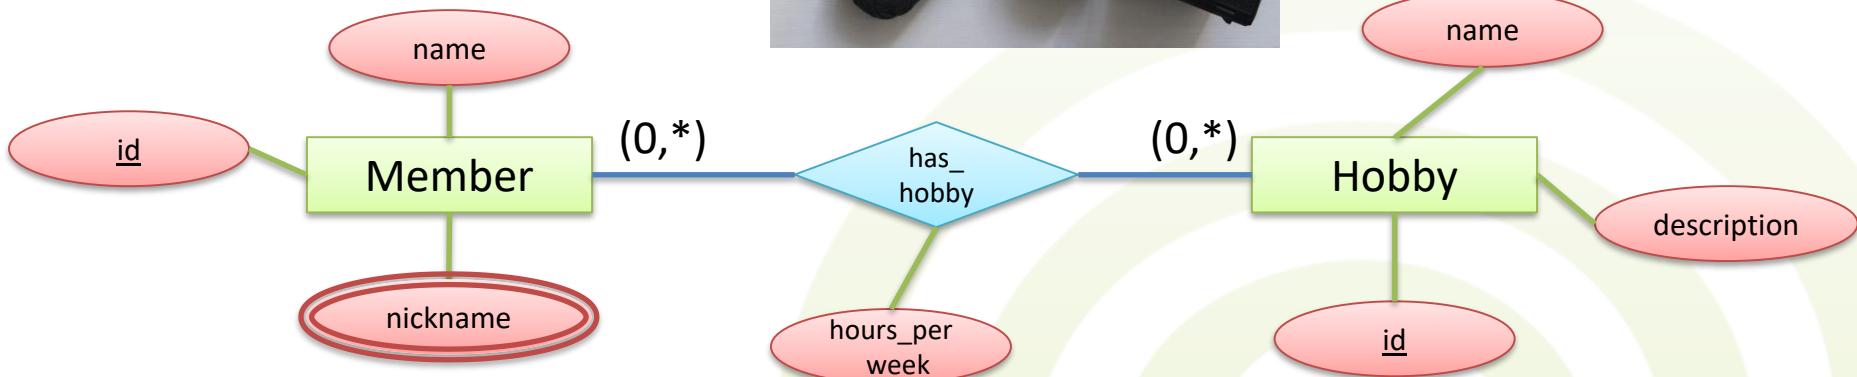
| artist_id | album_id | album_name     |
|-----------|----------|----------------|
| 3         | 1        | Schizophrenic  |
| 4         | 2        | Eat the rich   |
| 3         | 3        | Crave (single) |



# 9.1 SQL DDL: Referential Integrity

## Example

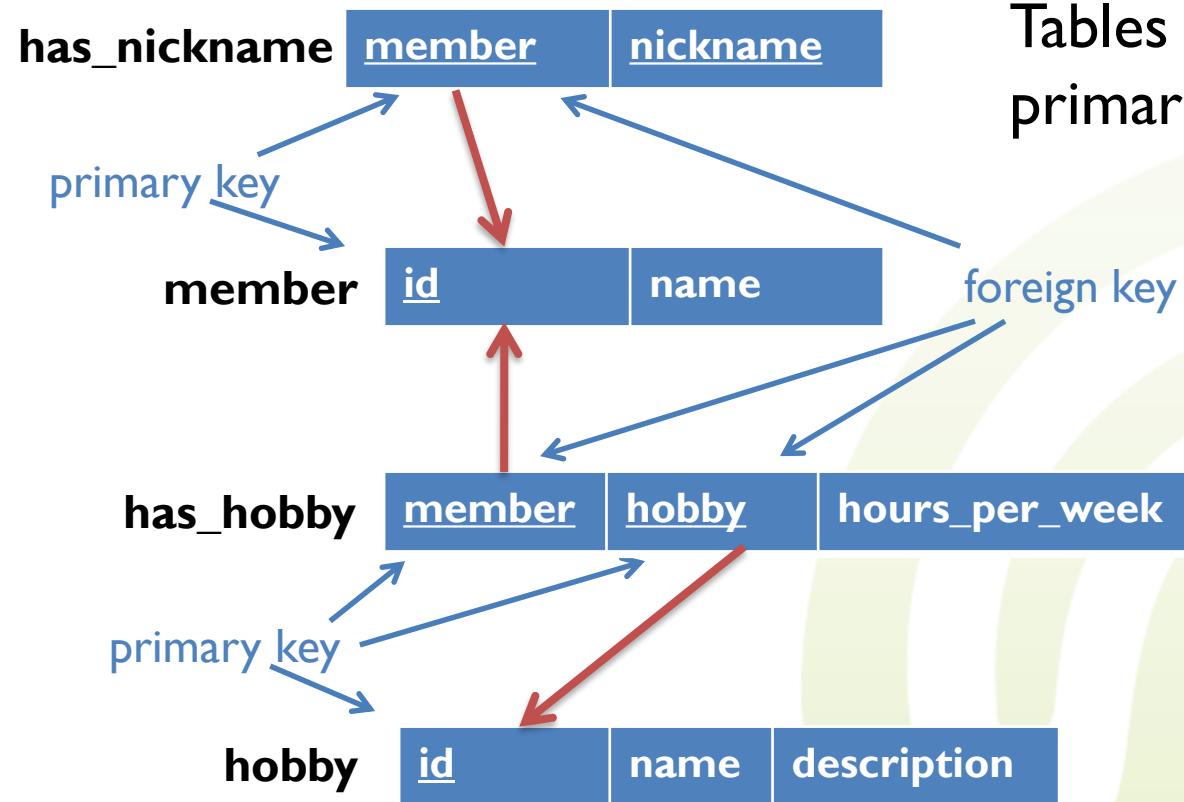
### Conceptual ER schema





# 9.1 SQL DDL: Referential Integrity

## Resulting tables



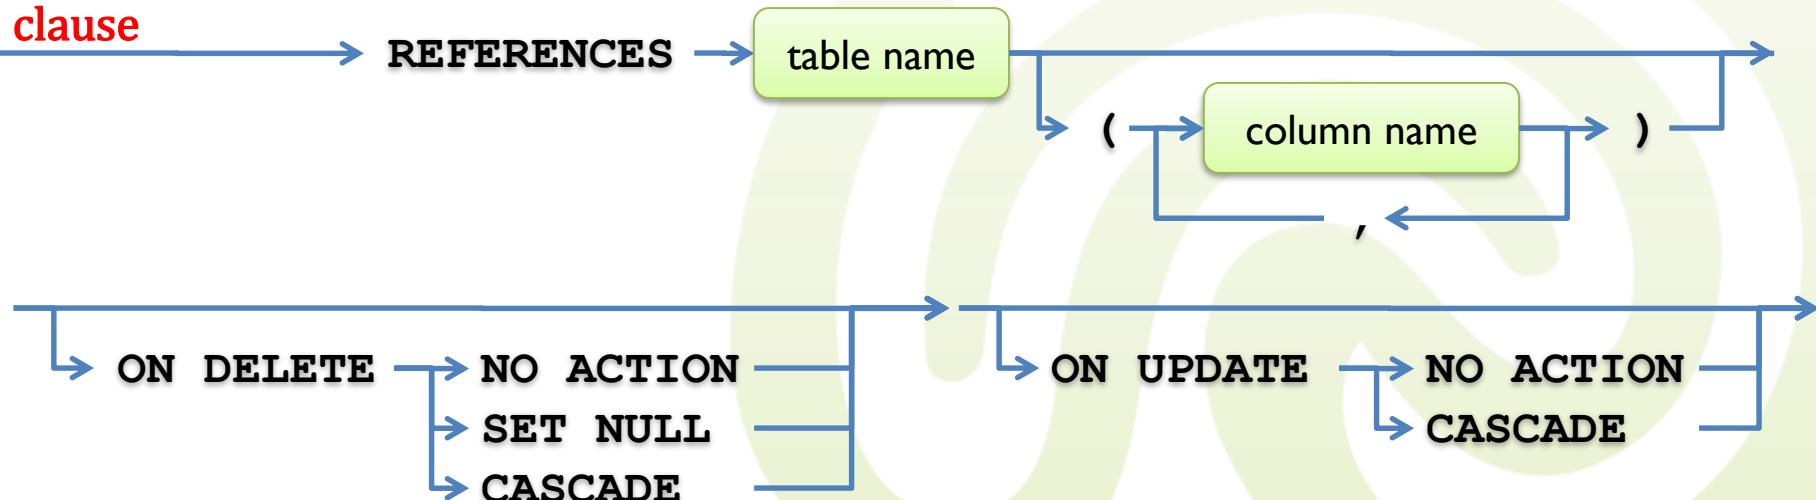
Tables refer to others by primary keys and foreign keys



## 9.1 SQL DDL: Referential Integrity

- Referential integrity can be defined using the **REFERENCES clause**
  - either used by constraints in **column options** or within **table constraints**
  - if no attribute is selected, it refers to the primary key

**REFERENCES-  
clause**





## 9.1 SQL DDL: Referential Integrity

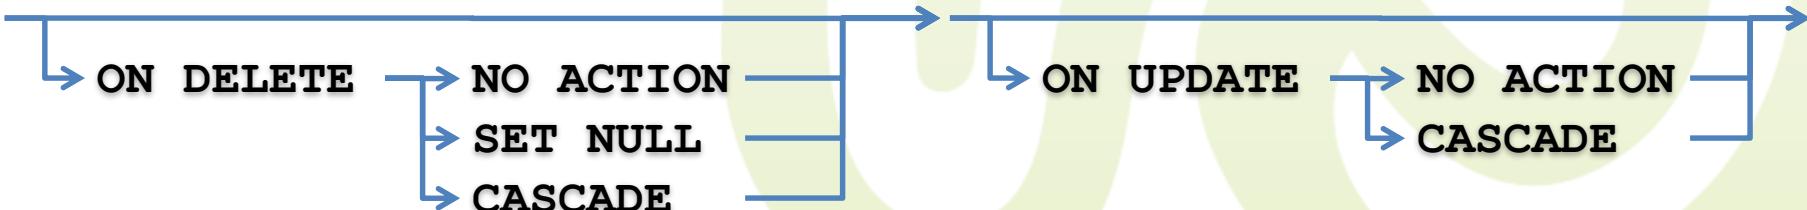
- Example

- `CREATE TABLE employee(`  
    `id INTEGER NOT NULL PRIMARY KEY,`  
    `name VARCHAR(100) NOT NULL`  
`)`
  - `CREATE TABLE managed_by(`  
    `employee INTEGER NOT NULL`  
        `REFERENCES employee,`  
    `manager INTEGER NOT NULL`  
        `REFERENCES employee`  
`)`



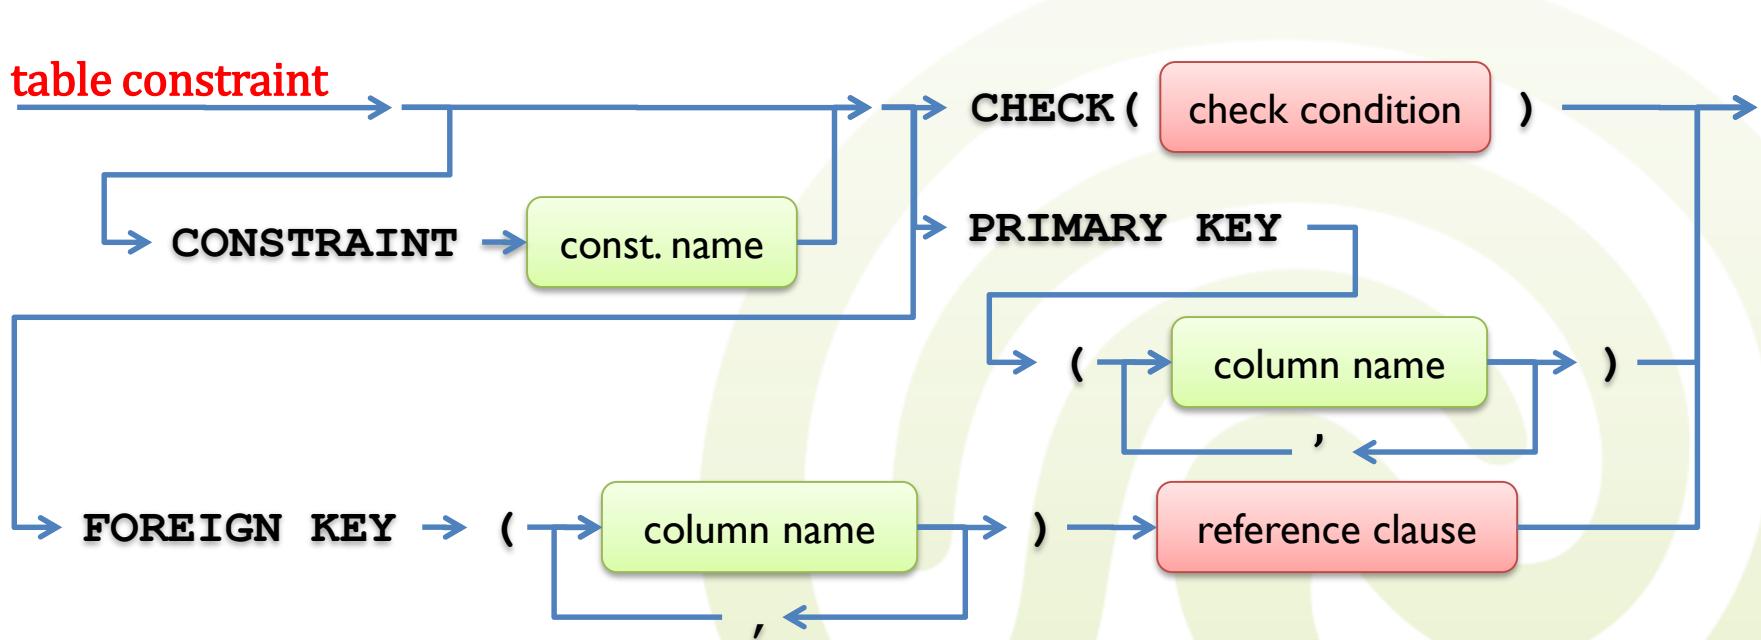
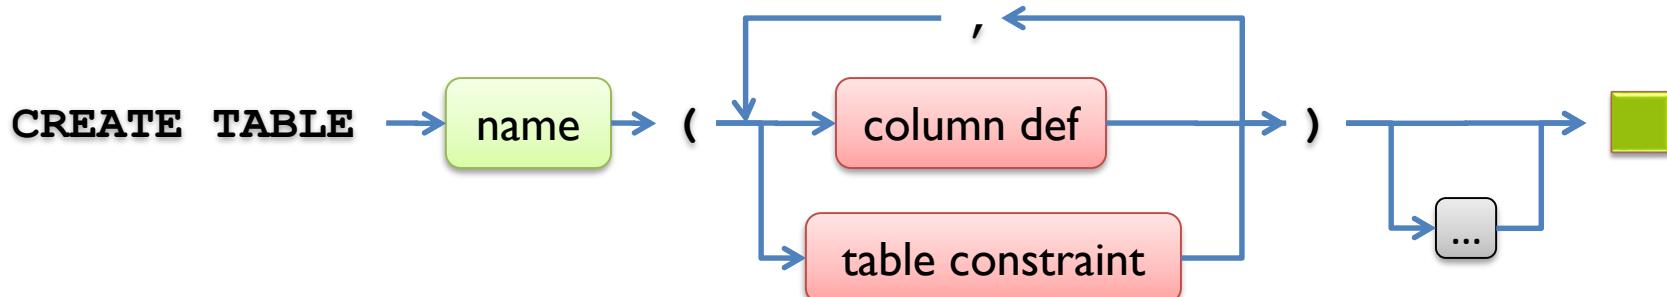
## 9.1 SQL DDL: Referential Integrity

- Optionally, you may specify what happens if a row that is referenced will be deleted or modified
  - **ON DELETE**: if a referenced row is deleted, ...
    - **NO ACTION**: ...reject the deletion (that is, it cannot be performed)
    - **SET NULL**: ...delete it and set all referencing foreign keys to **NULL**
    - **CASCADE**: ...delete it along with all rows referring to it
  - **ON UPDATE**: if the primary key of a referenced row is modified, ...
    - **NO ACTION**: ...reject the modification (that is, it cannot be performed)
    - **CASCADE**: ...change all values of referencing foreign keys
  - default
    - **ON DELETE NO ACTION ON UPDATE NO ACTION**





# 9.1 SQL DDL:Table Constraints



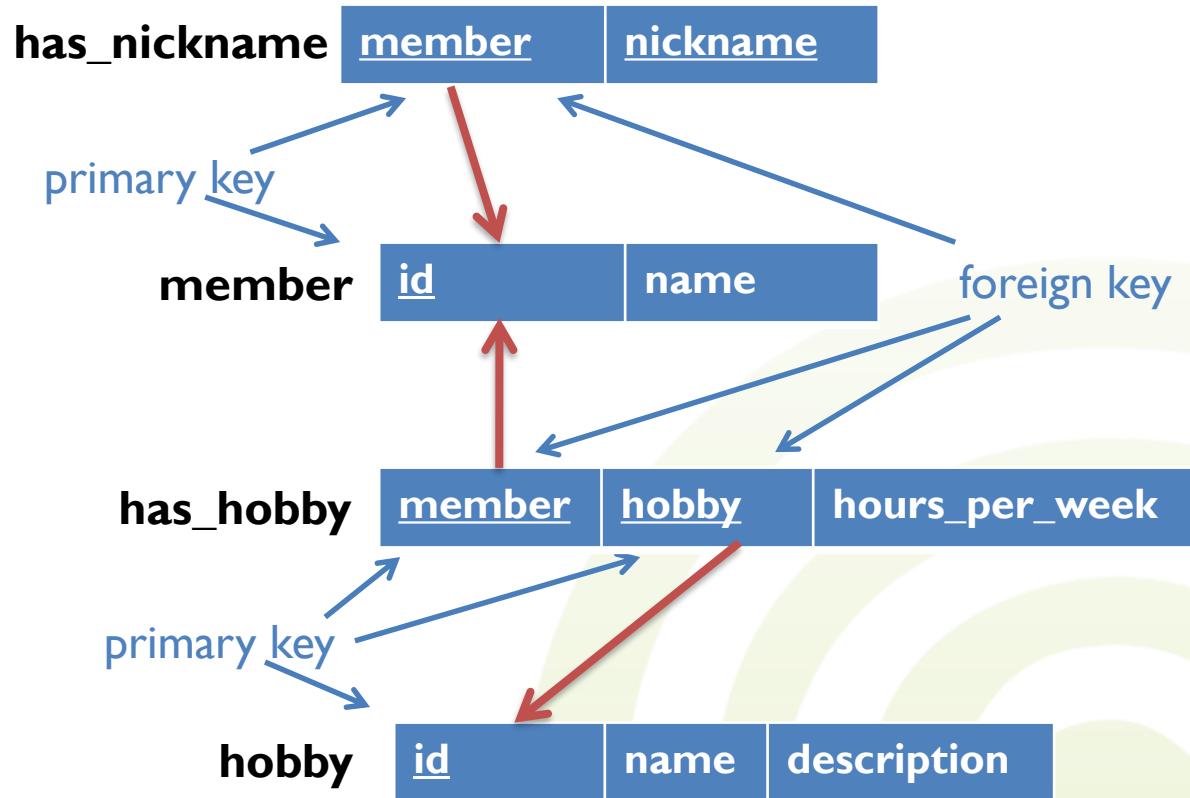


## 9.1 SQL DDL:Table Constraints

- Table constraints behave similar to constraints in column options
  - if no name is provided, a **name is automatically generated**
  - the **CHECK condition** may contain any Boolean predicate
  - in contrast to column options, table constraints may declare **primary keys** consisting of **multiple attributes**
  - **foreign keys** declare references to primary keys of other tables
    - see referential integrity



# 9.1 SQL DDL: Table Example





# 9.1 SQL DDL:Table Example

```
CREATE TABLE member (
    id INTEGER NOT NULL PRIMARY KEY,
    name VARCHAR(100)
);

CREATE TABLE hobby(
    id INTEGER NOT NULL PRIMARY KEY,
    name VARCHAR(100),
    description VARCHAR(255)
);

CREATE TABLE has_nickname (
    member INTEGER REFERENCES member ← link has_nickname to member
        ON DELETE CASCADE ← delete nickname if member is deleted
        ON UPDATE CASCADE, ← update nickname if member is updated
    nickname VARCHAR(100) NOT NULL,
    PRIMARY KEY (member, nickname) ← composed primary key
);
```



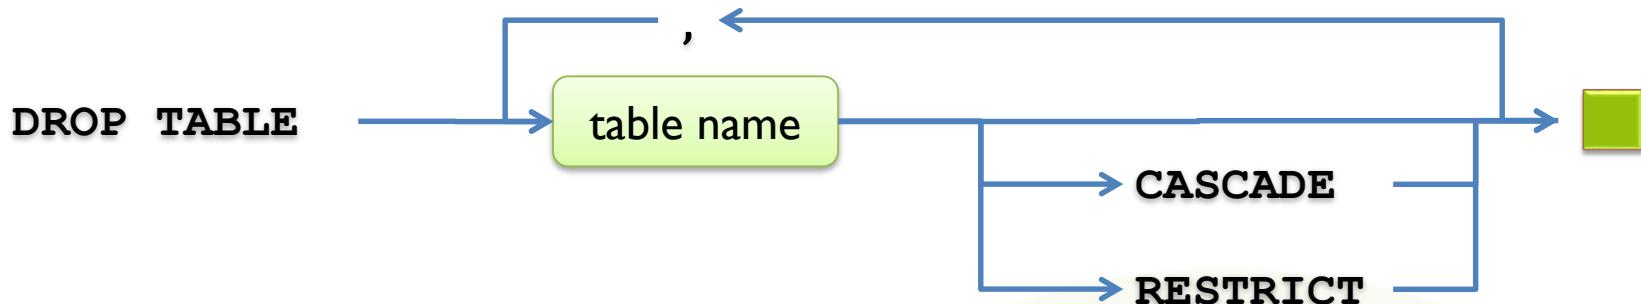
# 9.1 SQL DDL: Table Example

- `CREATE TABLE has_hobby (`  
    `member INTEGER NOT NULL,`  
    `hobby INTEGER NOT NULL,`  
    `hours_per_week INTEGER NOT NULL,`  
    `PRIMARY KEY (member, hobby),`  
    `FOREIGN KEY (member) REFERENCES member`  
        `ON DELETE CASCADE`  
        `ON UPDATE CASCADE,`  
    `FOREIGN KEY (hobby) REFERENCES hobby`  
        `ON DELETE CASCADE`  
        `ON UPDATE CASCADE`  
)



# 9.1 SQL DDL: Drop Tables

- For **deleting** tables, there is the **DROP TABLE** command

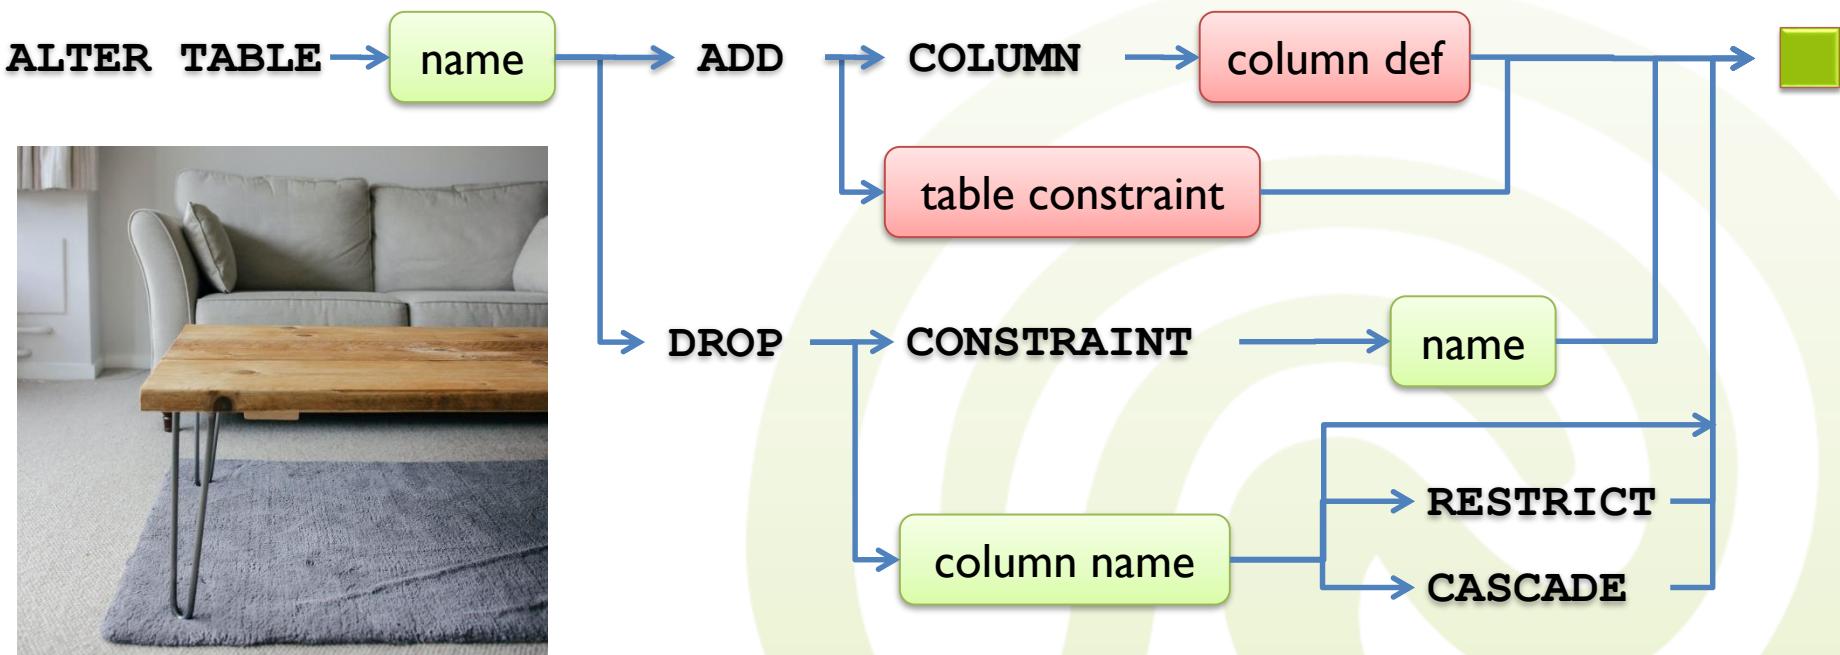


- if **RESTRICT** is used, you may only drop empty tables that are not referenced by any other table
- if **CASCADE** is used, all referencing tables are also deleted (including all stored rows)
- if neither is used, the table does not have to be empty, but must not be referenced by another one
- example
  - DROP TABLE member CASCADE, hobby CASCADE**



# 9.1 SQL DDL: Alter Tables

- After a table has been created, you may **alter** it by adding/removing **columns** or **constraints**





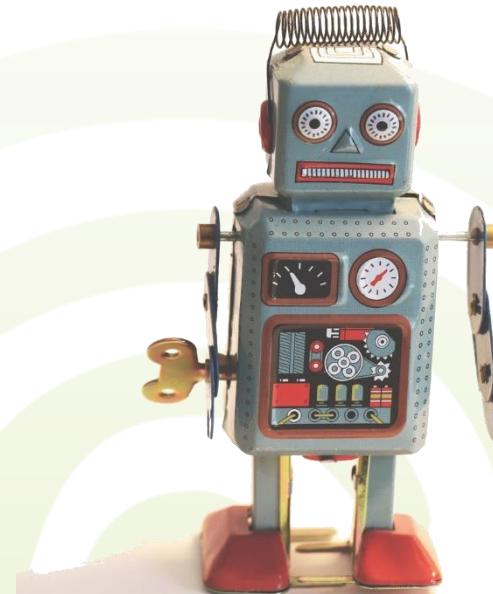
# 9.1 SQL DDL: Alter Tables

- if you add a **new column** with a **NOT NULL** constraint, you also need to provide a **default value**
- when **dropping a column**, you must either choose
  - **CASCADE** to also delete any views, indexes, and constraints dependent on that column
  - **RESTRICT** to allow the drop only if there is no referring column (default)
- if the **name** of a constraint is **auto-generated**, you need to look it up in the **system catalog**
- example
  - **ALTER TABLE** has\_hobby **DROP** **hours\_per\_week**  
**ALTER TABLE** has\_hobby  
**ADD COLUMN** **since** **DATE**



# 9 SQL 2

- SQL data definition language
- **SQL data manipulation language  
(apart from SELECT)**
- SQL  $\neq$  SQL
- Some advanced SQL concepts





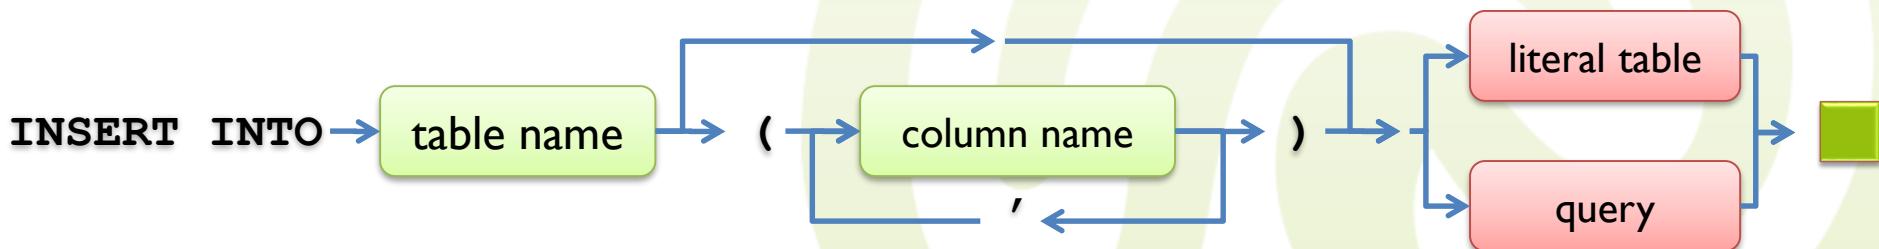
## 9.2 SQL DML

- **Data definition language (DDL)**
  - creating, changing, altering schemas, tables, ...
    - CREATE SCHEMA
    - CREATE TABLE
    - ALTER TABLE
    - DROP TABLE
- **Data manipulation language (DML)**
  - querying
    - SELECT
  - adding and updating data
    - INSERT INTO
    - UPDATE
    - DELETE



## 9.2 SQL DML: Insert

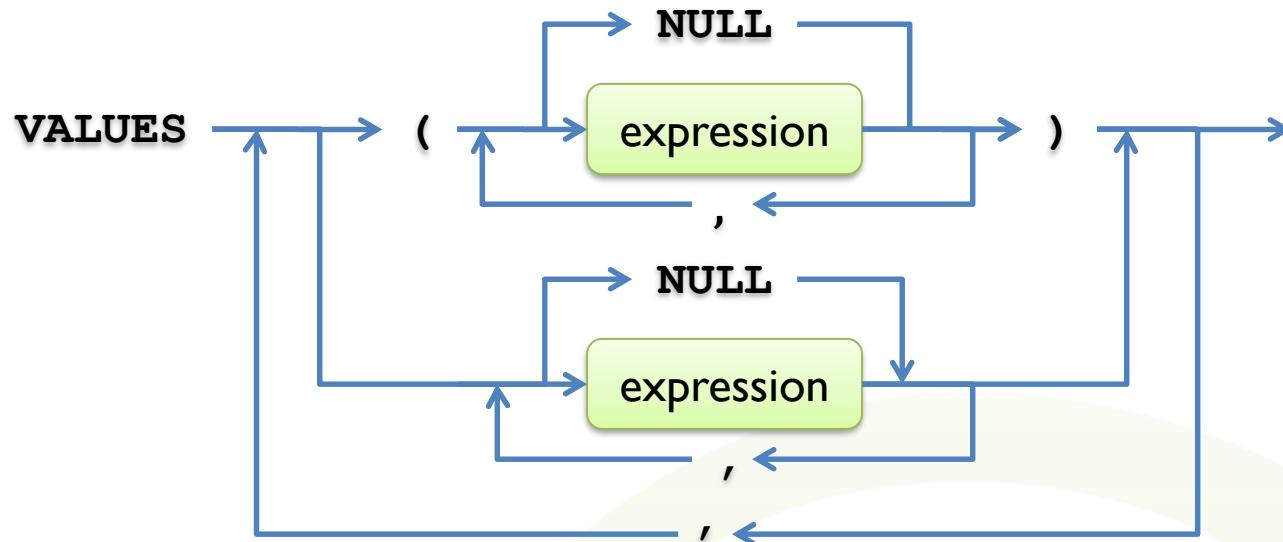
- Now we have wonderful, empty tables
- We need to put data into them!
  - **INSERT INTO** statement
  - you can specify into what **columns** you want to insert data
    - default: all columns
  - new values are stated as a **literal table or inline view (query)**
    - of course the **attribute domains** have to match





## 9.2 SQL DML: Insert

literal table



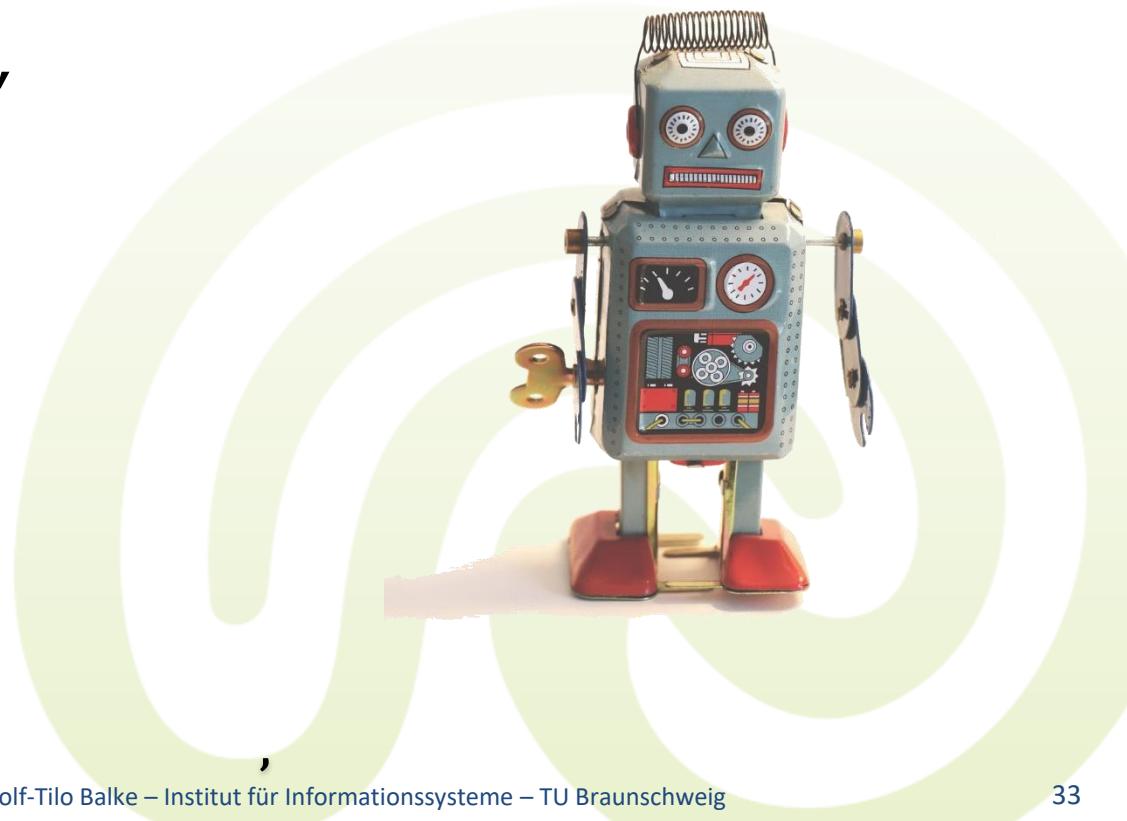
- A **literal table** is defined extensionally:

|                                                               |                                                                                                             |          |             |          |         |
|---------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|----------|-------------|----------|---------|
| <code>VALUES ('Artemis', 'Towel')</code>                      | <table border="1"><tr><td>Artemis</td><td>Towel</td></tr></table>                                           | Artemis  | Towel       |          |         |
| Artemis                                                       | Towel                                                                                                       |          |             |          |         |
| <code>VALUES ('Karl', 'Sparks'), ('Hannah', 'Antenna')</code> | <table border="1"><tr><td>Karl</td><td>Sparks</td></tr><tr><td>Hannah</td><td>Antenna</td></tr></table>     | Karl     | Sparks      | Hannah   | Antenna |
| Karl                                                          | Sparks                                                                                                      |          |             |          |         |
| Hannah                                                        | Antenna                                                                                                     |          |             |          |         |
| <code>VALUES 'Sokrates', ('Professor Y'), 'Asteriks'</code>   | <table border="1"><tr><td>Sokrates</td></tr><tr><td>Professor Y</td></tr><tr><td>Asteriks</td></tr></table> | Sokrates | Professor Y | Asteriks |         |
| Sokrates                                                      |                                                                                                             |          |             |          |         |
| Professor Y                                                   |                                                                                                             |          |             |          |         |
| Asteriks                                                      |                                                                                                             |          |             |          |         |



## 9.2 SQL DML: Insert

- **INSERT INTO** member(id, real\_name) **VALUES**  
**(1, 'Sammy Riegel'),**  
**(2, 'Morris Hinz')**  
**INSERT INTO** has\_nickname **VALUES**  
**(1, 'FCG'),**  
**(1, 'Letters'),**  
**(2, 'Moritz'),**  
**(2, 'Boris')**





## 9.2 SQL DML: Insert

- Of course, subqueries may also be used in **INSERT** statements

```
— INSERT INTO members_starting_with_a(  
    SELECT * FROM member  
    WHERE name LIKE 'A%'  
)
```





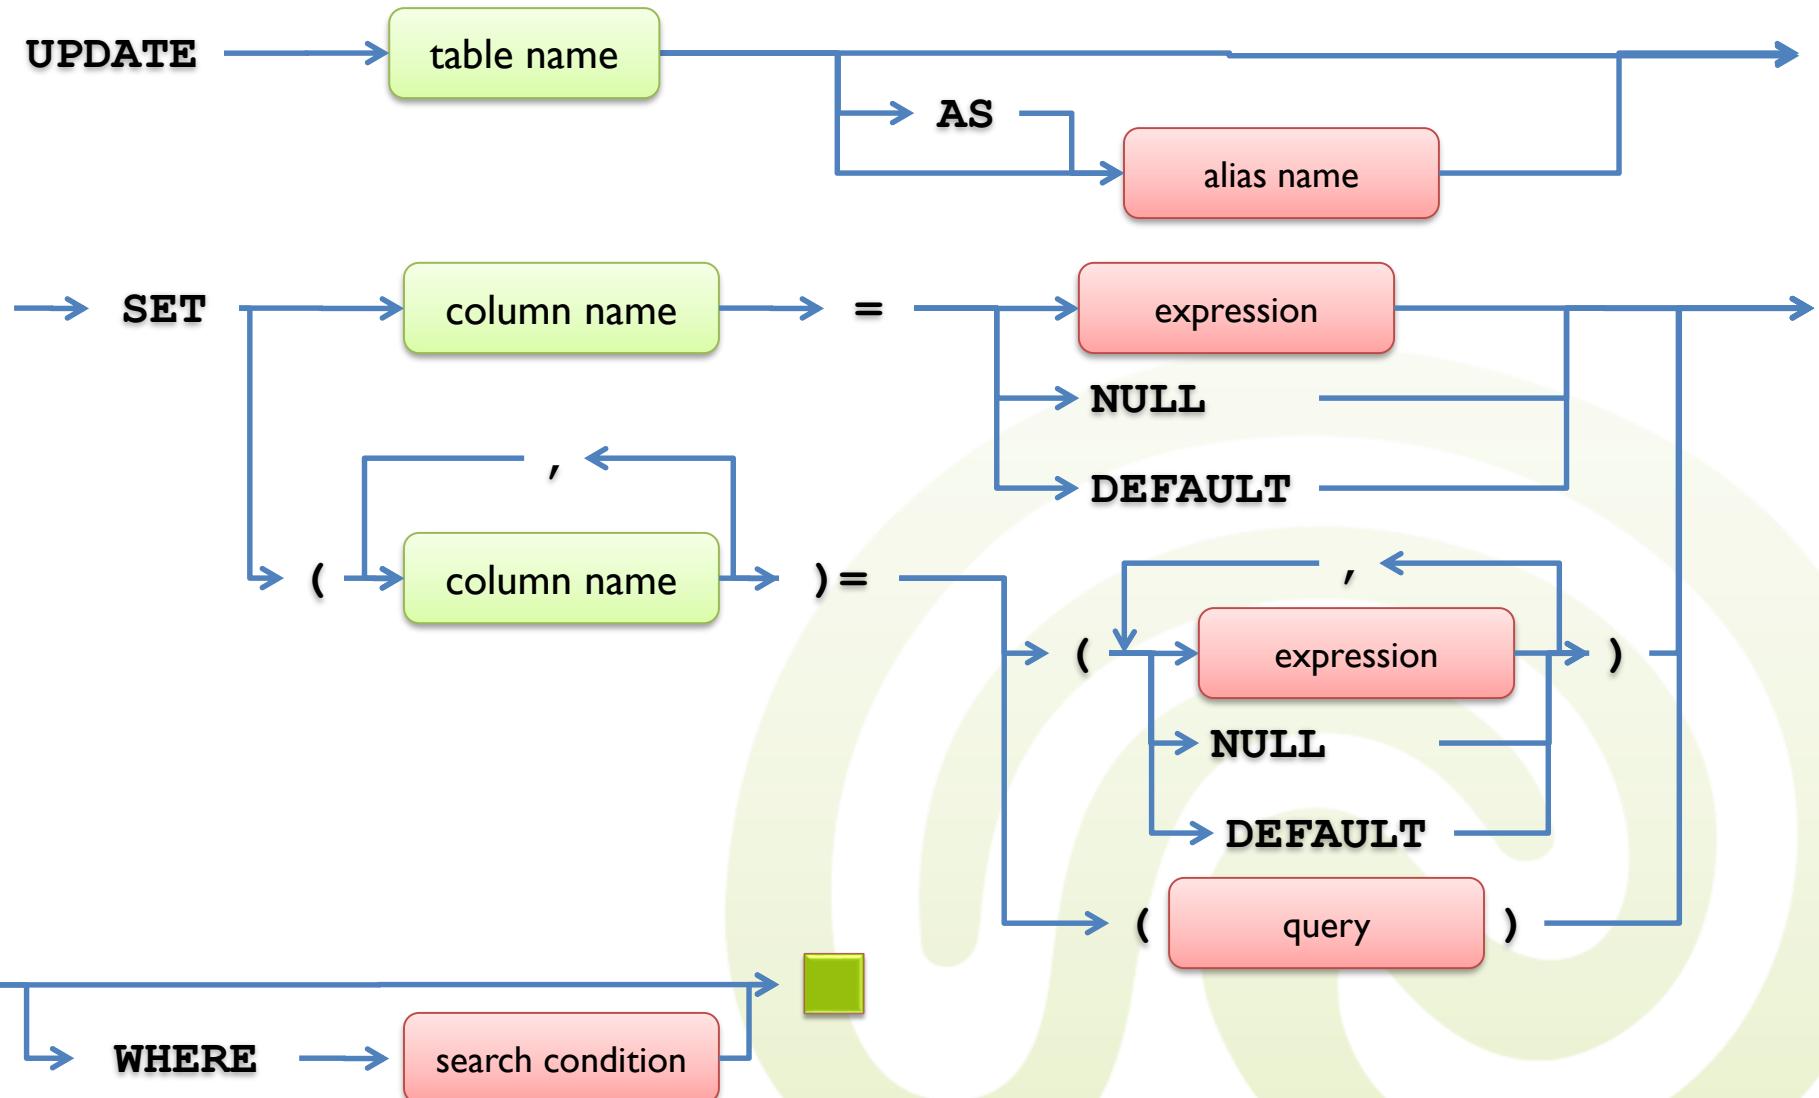
## 9.2 SQL DML: Update

- Existing rows can also be changed using the **UPDATE statement**
  - very similar to the **SELECT** statement
  - update **finds rows** fulfilling a **given condition** and **changes** some of its **rows** by assigning **new values**





# 9.2 SQL DML: Update





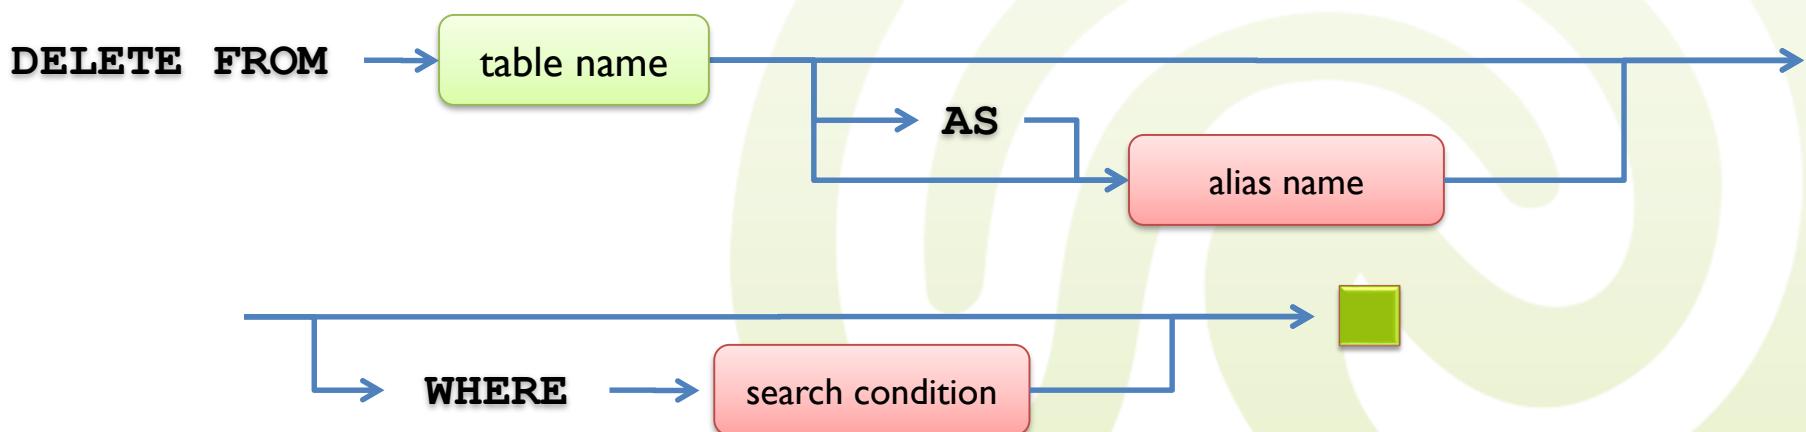
## 9.2 SQL DML: Update

- Replace the name of each member with **NULL**
  - `UPDATE member SET name = NULL`
- Multiply all `hours_per_week` values by 10
  - `UPDATE has_hobby SET hours_per_week = hours_per_week * 10`
- Change the name of the member with id 1
  - `UPDATE member SET name= 'John of Warrington' WHERE id = 1`
- Change name and id of Rick Chances
  - `UPDATE member SET (id, name) = ('3', 'Pickle Rick') WHERE name = 'Rick Chances'`
  - Change of id is propagated to other tables when `ON UPDATE CASCADE` is used in table definition
- Again, subqueries can be used in the **WHERE** clause



## 9.2 SQL DML: Delete

- The **DELETE statement** is used to delete rows from a table
  - deletes all rows satisfying a certain search condition
  - example
    - delete Pickle Rick
      - `DELETE FROM member WHERE name = 'Pickle Rick'`
    - delete all members
      - `DELETE FROM member`





## 9.2 SQL DML: Delete

- Again, subqueries can be used here
  - `DELETE FROM member m  
WHERE NOT EXISTS (  
 SELECT * FROM has_hobby h  
 WHERE h.member = m.id  
)`





# 9 SQL 2

- SQL data definition language
- SQL data manipulation language  
(apart from SELECT)
- **SQL ≠ SQL**
- Some advanced SQL concepts





# 9.3 SQL ≠ SQL

*Detour*

- First, the **good news**:
  - SQL has been standardized by the ISO in 1987
  - the standard is well-maintained and under active development
    - SQL-86, SQL-89, **SQL-92**, **SQL:1999**, **SQL:2003**, SQL:2006, SQL:2008, SQL:2011, SQL:2016, SQL:2019, SQL:2023
  - many *big* database vendors participate in the standardization process
    - IBM, Postgres, Microsoft, Oracle, Sybase, ...





## 9.3 SQL ≠ SQL

*Detour*

### A timeline of SQL standardization:

- 1986
  - ANSI SQL
  - relations, attributes, views
  - **SELECT ... FROM ... WHERE ...**
- 1987
  - SQL-86 (ISO/IEC 9075:1986)
- 1989
  - SQL-89 (SQLI)
    - ≈ SQL-86 + restricted referential integrity





## 9.3 SQL ≠ SQL

*Detour*

- 1992
  - SQL-92 (SQL2)
  - 3 parts, 1120 pages
  - Entry Level
    - ≈ SQL-89 + **CHECK** (attribute)
  - Intermediate Level
    - ⊇ Entry Level + domains, **CHECK** (relation),  
**CASE**, **CAST**, **JOIN**, **EXCEPT**, **INTERSECT**
  - Full Level
    - ⊇ Intermediate Level + assertions, nested select,  
nested from



# 9.3 SQL ≠ SQL

*Detour*

- 1999/2000
  - SQL:1999 (SQL3)
  - 5 parts, 2084 pages
  - ≈ SQL-92 + object-orientation, **recursive queries**, triggers, OLAP, user-defined types, regular expressions
    - Boolean data type
  - computationally complete, object-oriented database programming language, descriptive and procedural
  - core (about 180 features)
    - ≈ SQL92 Entry Level + parts of Intermediate and Full Level
  - 9 Packages (about 240 features)
    - enhanced date/time, enhanced integrity, OLAP, PSM, CLI, basic object support, enhanced object support, trigger, SQL/MM



# 9.3 SQL ≠ SQL

*Detour*

- Recursive queries (SQL:1999):
  - *How to find all rivers that flow into the North Sea?*
  - Multiple joins?



| flows_into | river  | mouth     |
|------------|--------|-----------|
|            | Oker   | Aller     |
|            | Aller  | Weser     |
|            | Weser  | North Sea |
|            | Elbe   | North Sea |
|            | Edder  | Wietze    |
|            | Flöth  | Wietze    |
|            | Wietze | Aller     |
|            | Isar   | Danube    |
|            | Inn    | Danube    |



# 9.3 SQL ≠ SQL

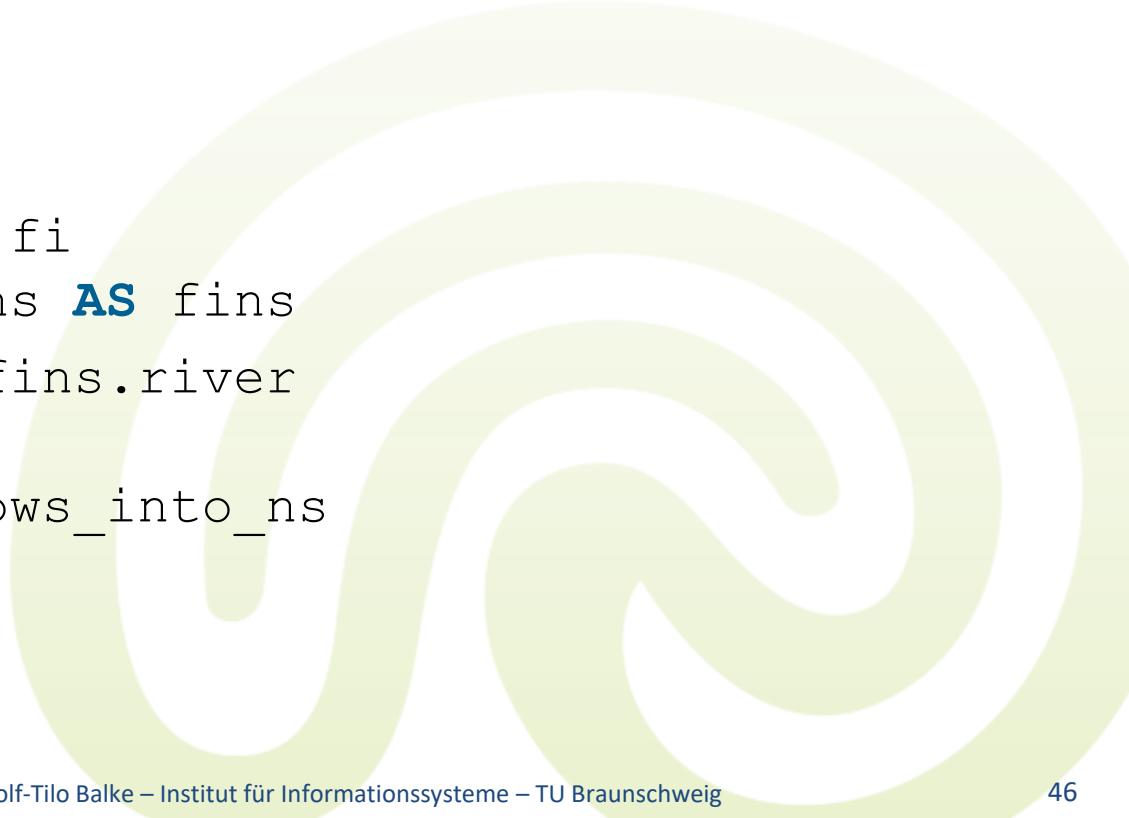
*Detour*

- Solution

```
WITH RECURSIVE flows_into_ns(river) AS (
    SELECT river
    FROM flows_into
    WHERE mouth = 'North Sea'

    UNION

    SELECT fi.river
    FROM flows_into AS fi
    JOIN flows_into_ns AS fins
        ON fi.mouth = fins.river
)
SELECT river FROM flows_into_ns
```





# 9.3 SQL ≠ SQL

*Detour*

- 2003
  - SQL:2003
  - 14 parts, 3606 pages
  - **MULTISET** as an explicit construct (with numerous operations, such as: **MULTISET UNION, MULTISET EXCEPT, MULTISET INTERSECT, CARDINALITY**)
  - sequence generators
    - **CREATE SEQUENCE <sequence name> AS <type name>**  
**[START WITH <value>]** **[INCREMENT BY <value>]**  
**[NO MINVALUE | MINVALUE <value>]**  
**[NO MAXVALUE | MAXVALUE <value>]**  
**[NO CYCLE | CYCLE]**
    - Many databases use non-standard syntax for this task...
  - base type **XML** for mappings between SQL and XML
  - **CREATE TABLE AS / CREATE TABLE LIKE**



# 9.3 SQL ≠ SQL

*Detour*

- 2006
  - SQL:2006
  - successor of SQL:2003
  - a new extension for XML handling
    - importing, storing, querying, and manipulating XML data
    - support for XQuery
    - concurrently access (object-)relational data and XML data
- 2008
  - SQL:2008
  - successor of SQL:2006
  - *maintenance release* (some new but minor features)



# 9.3 SQL ≠ SQL

*Detour*

- 2011
  - SQL:2011
  - successor of SQL:2008
  - adds support for **time periods**



```
CREATE TABLE mutation (
    name VARCHAR(255) PRIMARY KEY NOT NULL,
    person INTEGER REFERENCES person(id),
    from DATE,
    until DATE,
    PERIOD FOR period_of_validity (from, until));
SELECT name, person FROM mutation
WHERE period_of_validity CONTAINS '2014-12-19';
```

- temporal predicates, temporal primary keys, temporal referential integrity
  - This might be very useful for work planning or scheduling applications



## 9.3 SQL ≠ SQL

*Detour*

- Well, here are the **bad news**
  - there are still **too many variants** of SQL  
**(both syntactic and semantic differences)**
    - true application portability remains a challenge
  - the standard has been used to introduce  
**two kinds of features**
    1. features that are well-understood and widely implemented
    2. new and largely untried technologies, hoping that vendors follow the lead and deliver new functionalities
  - **vendors don't care** too much about the standard



## 9.3 SQL ≠ SQL

*Detour*

- A **common myth** among software developers

If your application  
uses only  
standard SQL,  
then it is portable.

- If you don't believe me, here are some examples ...



## 9.3 SQL ≠ SQL

*Detour*

- ```
CREATE TABLE name (
    first VARCHAR(100),
    middle VARCHAR(100),
    last VARCHAR(100)
)
```

```
INSERT INTO name VALUES ('George', 'Walker', 'Bush')
INSERT INTO name VALUES ('Horst', '', 'Kr')
INSERT INTO name VALUES ('Angela', NULL, 'Merkel')
```

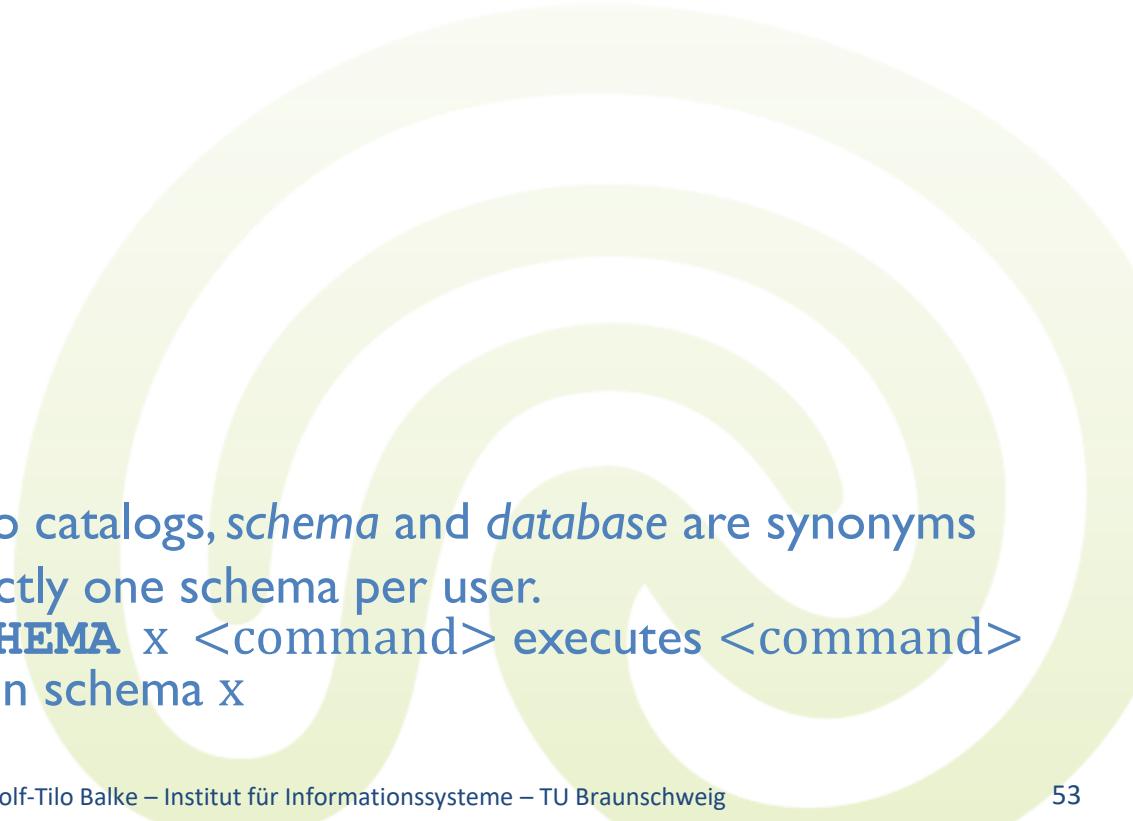
- ' '(empty string) means that we know that there is no middle name
- **NULL** means that we don't know whether there is a middle name
- Sounds like a good design? What do you think?
  - according to the SQL standard, this approach is fine ...
  - ... unless your RDBMS is Oracle (' ' is the same as **NULL**)



# 9.3 SQL ≠ SQL

*Detour*

- What about terminology?
  - the SQL standard defines the following notions
    - Environment
    - Cluster
    - Catalog
    - Schema
  - the reality
    - Database server
    - (unsupported)
    - Database
    - Schema
  - but attention
    - in MySQL, there are no catalogs, *schema* and *database* are synonyms
    - in Oracle, there is exactly one schema per user.  
**CREATE/ALTER SCHEMA** x <command> executes <command> on all objects located in schema x





## 9.3 SQL ≠ SQL

*Detour*

- The **statement terminator** ;
  - according to the SQL standard, (almost) every SQL statement has to be terminated by a semicolon
  - *What's happening in practice?*
    - many RDBMS treat the terminator as being optional (which is fine, but may cause some problems)
      - PostgreSQL does not require semicolons for single commands, but multiple commands must be separated by semicolons
    - some RDBMS either strictly require a terminator or complain if it is present
    - in some RDBMS, this behavior can be configured ...
  - summary:  
**No matter what you do, it causes problems!**



# 9.3 SQL ≠ SQL

*Detour*

- The **BOOLEAN** data type

- ```
CREATE TABLE customers (
    id INTEGER PRIMARY KEY,
    name VARCHAR(100),
    is_vip BOOLEAN,
    is_blacklisted BOOLEAN
)
```



```
SELECT id, name FROM customers
WHERE is_vip AND NOT is_blacklisted
```

- practice?

- not supported by DB2, and MS SQL Server
  - official workarounds: use CHAR or INTEGER ...
- supported by Oracle, MySQL and PostgreSQL
  - where in MySQL BOOLEAN is just a short hand for TINYINT(1) ...



# 9.3 SQL ≠ SQL

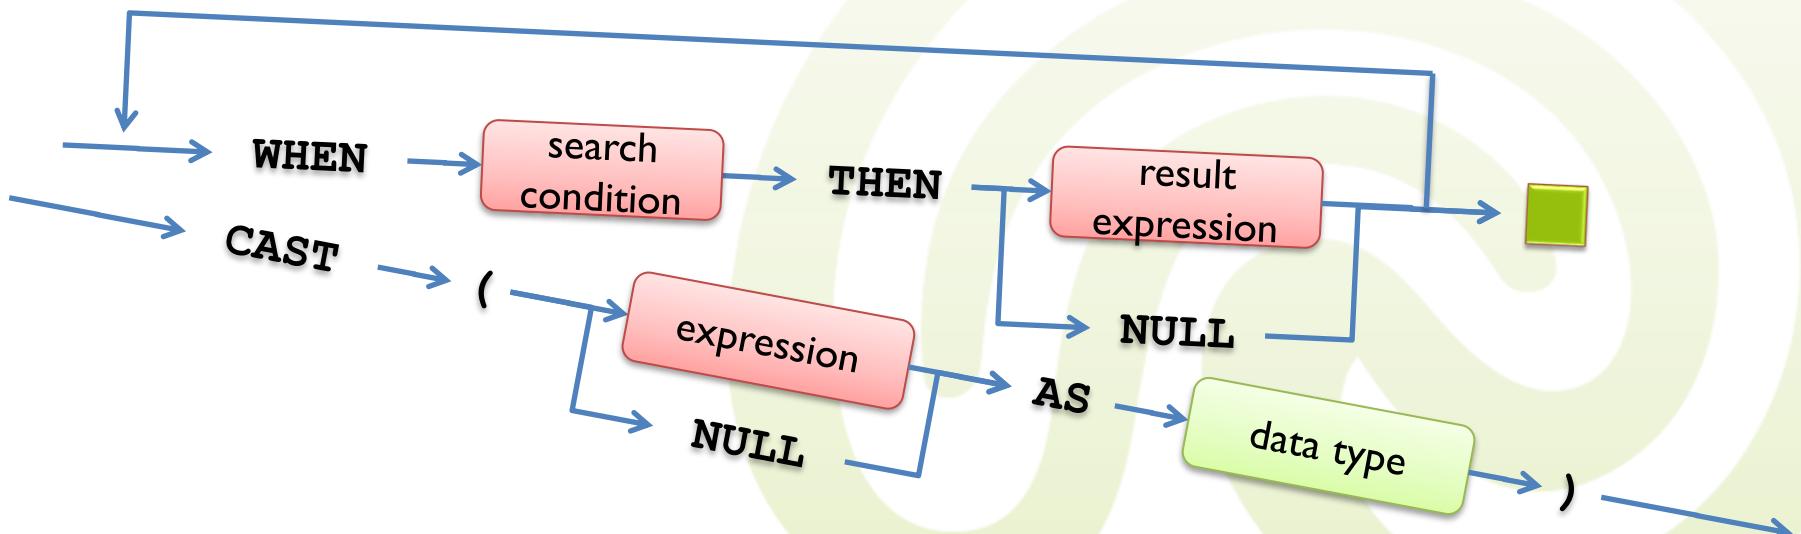
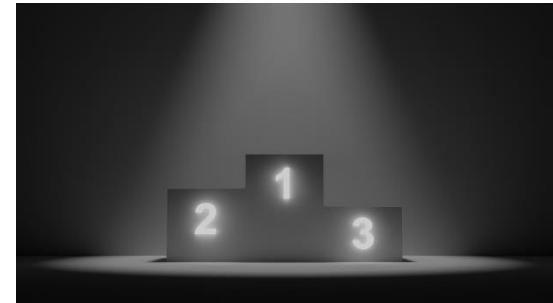
*Detour*

- **Summary**
  - **SQL is not SQL**
  - in some cases, even **identical SQL** statements **work differently** on different RDBMS
- **Current trends?**
  - open-source RDBMS (PostgreSQL, MySQL, Derby, ...) typically try to adhere to the standard
    - however, many advanced features are not supported yet
    - new features in PostgreSQL 14:
      - parallel query execution on remote databases
      - query pipelining: send multiple queries without waiting for responses
      - ...



# 9 SQL 2

- SQL data definition language
- SQL data manipulation language  
(apart from **SELECT**)
- $\text{SQL} \neq \text{SQL}$
- **Some advanced SQL concepts**

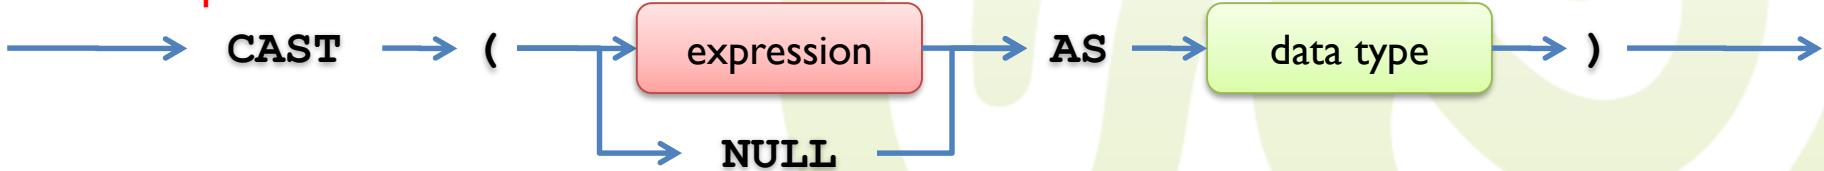




## 9.4 Type Casting

- SQL is a **strongly typed** language
  - Basically, this means that e.g. **INTEGER** is different from **VARCHAR(100)**
- If data types are incompatible, **type casting** may be used to make them compatible
  - **CAST** expression
  - during casting, precision may be lost (e.g. **FLOAT** → **INTEGER**)
  - example
    - **CAST** (hours\_per\_week **AS NUMERIC(3, 2)**)  
**CAST** (nickname || name **AS CHAR(255)**)
  - List possible type castings in **psql** with `\dCS <type>`

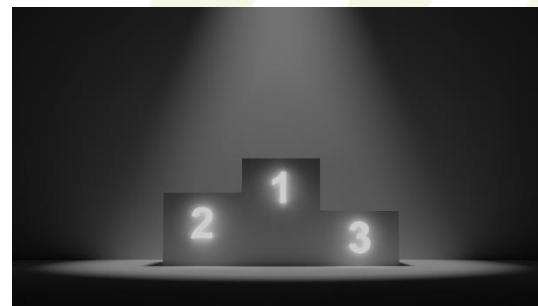
**CAST** expression





## 9.4 Ranking Functions

- Since SQL:2003, there are special functions for working with **result lists**
- Examples
  - output only every other row of the list
  - create a ranking with explicit ranks (1, 2, 3, ...)
  - on what rank position is some given row?





## 9.4 Ranking Functions

- **ROW\_NUMBER()** returns the **position** of each row in the result list
- Example

```
SELECT name, salary,  
       ROW_NUMBER() OVER (  
           ORDER BY salary DESC  
       ) AS pos  
FROM person
```

| person | name      | salary    |
|--------|-----------|-----------|
|        | Simon     | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Christoph | 45000     |



| name      | salary    | pos |
|-----------|-----------|-----|
| Larry     | 200000000 | 1   |
| Wolf-Tilo | 75000     | 2   |
| Christoph | 45000     | 3   |
| Simon     | 45000     | 4   |

Depending on the implementation, the last two rows may switch positions

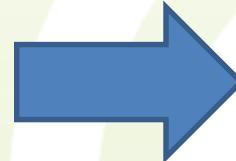


## 9.4 Ranking Functions

– example: At which position is Wolf-Tilo?

- **SELECT** name, salary,  
**ROW\_NUMBER()** **OVER** (  
    **ORDER BY** salary **DESC**  
    ) **AS** pos  
**FROM** person  
**WHERE** name = 'Wolf-Tilo'

| person | name      | salary    |
|--------|-----------|-----------|
|        | Simon     | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Christoph | 45000     |



| name      | salary | pos |
|-----------|--------|-----|
| Wolf-Tilo | 75000  | 2   |



## 9.4 Ranking Functions

– example: Show only rows at even positions.

```
• SELECT name, salary, ROW_NUMBER() OVER (
    ORDER BY salary DESC
) AS pos
FROM person
WHERE (pos % 2) = 0
```

modulo

| person | name      | salary    |
|--------|-----------|-----------|
|        | Simon     | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Christoph | 45000     |



| name      | salary | pos |
|-----------|--------|-----|
| Wolf-Tilo | 75000  | 2   |
| Simon     | 45000  | 4   |

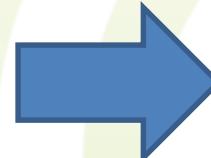


## 9.4 Ranking Functions

- **RANK()** returns the **rank** of each row in the result list
  - example

```
SELECT name, salary, RANK() OVER (
    ORDER BY salary DESC
) AS rank
FROM person
```

| person | name      | salary    |
|--------|-----------|-----------|
|        | Simon     | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Christoph | 45000     |



| name      | salary    | rank |
|-----------|-----------|------|
| Larry     | 200000000 | 1    |
| Wolf-Tilo | 75000     | 2    |
| Christoph | 45000     | 3    |
| Simon     | 45000     | 3    |

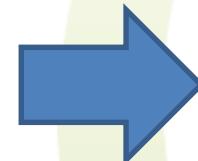


## 9.4 Ranking Functions

- **DENSE\_RANK()** works like **RANK()** but does not skip ranks on ties (as it is usually done)
  - example

```
SELECT name, salary, RANK() OVER (
    ORDER BY salary ASC
) AS rank, DENSE_RANK() OVER (
    ORDER BY salary ASC
) AS drank
FROM person
```

| person | name      | salary    |
|--------|-----------|-----------|
|        | Simon     | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Christoph | 45000     |



|  | name      | salary    | rank | drank |
|--|-----------|-----------|------|-------|
|  | Simon     | 45000     | 1    | 1     |
|  | Christoph | 45000     | 1    | 1     |
|  | Wolf-Tilo | 75000     | 3    | 2     |
|  | Larry     | 200000000 | 4    | 3     |



## 9.4 CASE Expressions

- Very often **codes** are used for storing more complex information
  - retrieving the account information for owner Shanks with appropriate account descriptions needs a join
  - Indicate all customers with a negative balance with the string *not creditworthy* in the query result

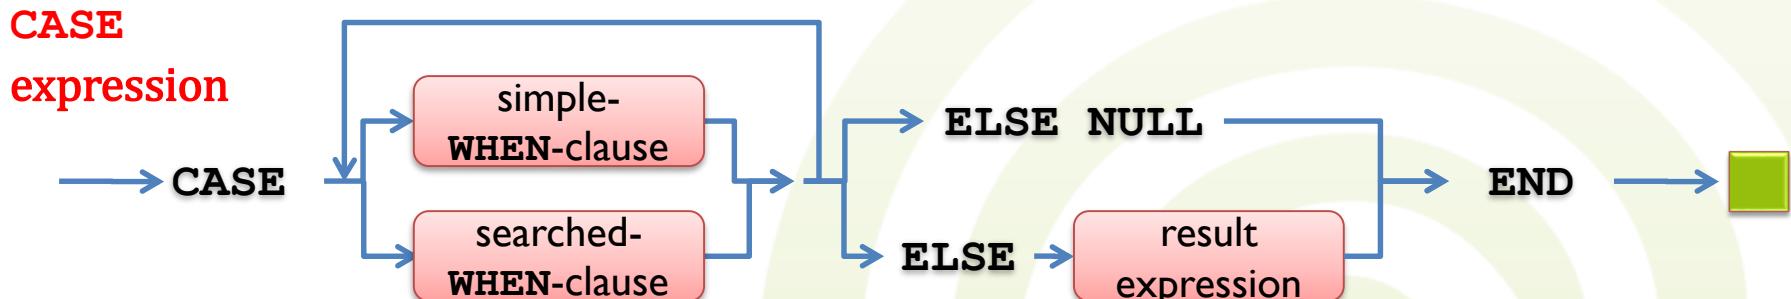
| account | owner  | balance  | type |
|---------|--------|----------|------|
|         | Shanks | 367,00   | 0    |
|         | Dido   | -675,00  | 0    |
|         | Shanks | 54987,00 | 1    |

| acc_type | type | description         |
|----------|------|---------------------|
|          | 0    | checking account    |
|          | 1    | savings account     |
|          | 2    | credit card account |



## 9.4 CASE Expressions

- The **CASE expression** allows a value to be selected based on the evaluation of one or more conditions (similar to *if-then-else*)
  - comes in two flavors



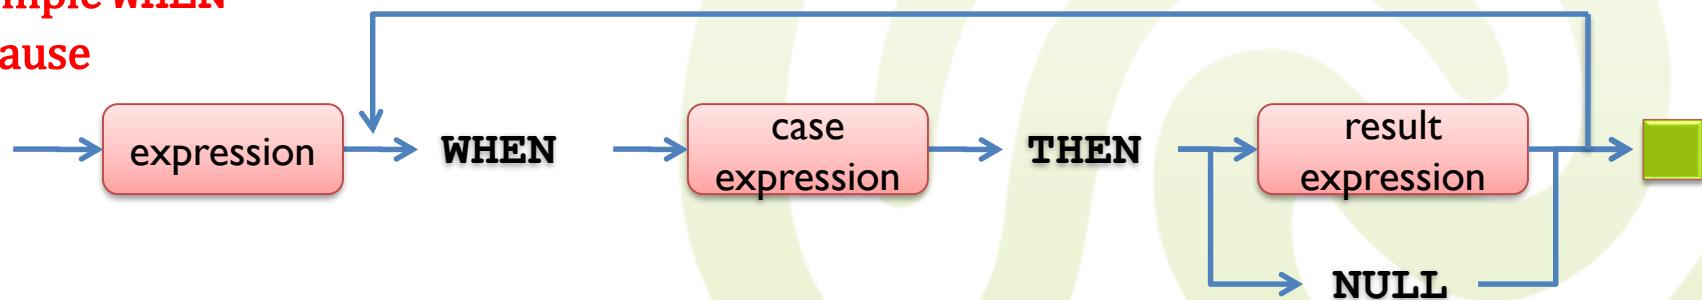


## 9.4 CASE Expressions

- the **simple WHEN clause**

- compares an expression to each case expression one by one
  - if expression is equal to search value, the corresponding result expression is returned
- if no match is found, then some default (**ELSE** clause) is returned
  - if **ELSE** is omitted, then **NULL** is returned

Simple WHEN clause





## 9.4 CASE Expressions

- **Example:** simple WHEN clause
  - directly decode the account type
    - **SELECT** owner,  
**CASE** type  
    **WHEN** 0 **THEN** 'checking account'  
    **WHEN** 1 **THEN** 'savings account'  
    **WHEN** 2 **THEN** 'credit card account'  
**END AS** verbose\_type  
**FROM** account

| account | owner    | balance | type |
|---------|----------|---------|------|
| Shanks  | 367,00   | 0       |      |
| Dido    | -675,00  | 0       |      |
| Shanks  | 54987,00 | 1       |      |



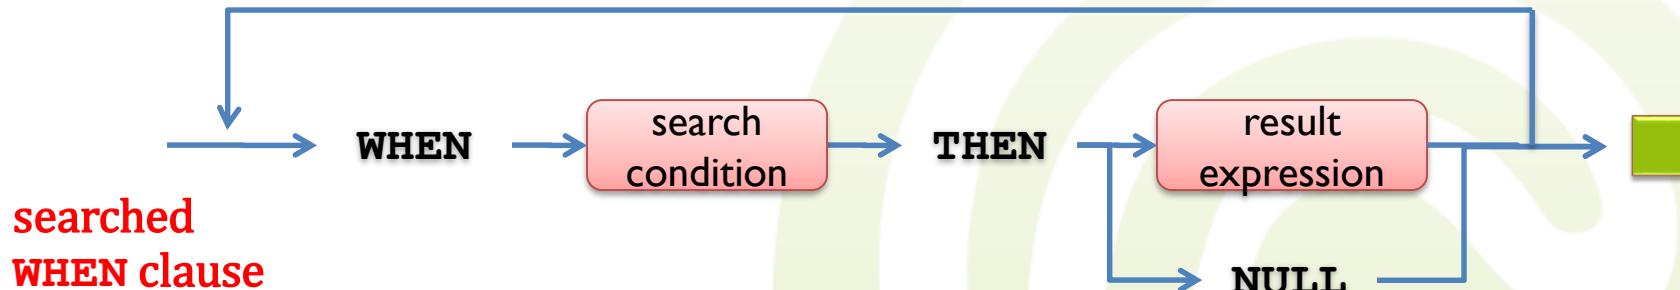
| owner  | verbose_type     |
|--------|------------------|
| Shanks | checking account |
| Dido   | checking account |
| Shanks | savings account  |



## 9.4 CASE Expressions

- the **searched WHEN clause**

- checks search conditions from left to right
- **stops as soon as a search condition evaluates to true**
  - returns the corresponding result then
- if **no condition is true**, the value given by the **ELSE clause** is returned (or **NULL**, if there is no **ELSE clause**)





## 9.4 CASE Expressions

- **Example:** searched **WHEN** clause

- retrieve credit rating of customers

- **SELECT** owner,

- CASE**

- WHEN** balance < 0 **THEN** 'not credit-worthy'

- WHEN** balance = 0 **THEN** 'questionable'

- ELSE** 'credit-worthy'

- END AS** credit\_worthyness

- FROM** account

- WHERE** type = 0

| account | owner  | balance  | type |
|---------|--------|----------|------|
|         | Shanks | 367,00   | 0    |
|         | Dido   | -675,00  | 0    |
|         | Shanks | 54987,00 | 1    |



| owner  | credit_worthyness |
|--------|-------------------|
| Shanks | credit-worthy     |
| Dido   | not credit-worthy |



# 9 More on SQL

- There are **many more SQL** statements than we are covering in our lectures
- Moreover, there are many different SQL dialects
- If you don't want to get mad, do the following
  - don't care too much about the SQL standard (unless you are actually implementing an RDBMS)
  - read the **SQL manuals** of your RDBMS!

**ORACLE**

**IBM** | **DB2**

**MySQL**



**Microsoft**  
**SQL Server**

Postgre**SQL**



# 9 More on SQL

- Example: PostgreSQL reference (version 14)
  - <https://www.postgresql.org/docs/14/reference.html>

The screenshot shows the PostgreSQL 14 documentation page for the `SELECT` command. The top navigation bar includes links for Home, About, Download, Documentation, Community, Developers, Support, Donate, and Your account. A search bar is also present. The main content area features a dark blue header with the date "10th November 2022: PostgreSQL 15.1, 14.6, 13.9, 12.13, 11.18, and 10.23 Released!" and a search bar. Below this, the `SELECT` command is highlighted in red, with its definition: "SQL Commands". Navigation links for Prev, Up, Home, and Next are available. The `Synopsis` section contains the detailed syntax of the `SELECT` command, which is a large block of text starting with `[ WITH [ RECURSIVE ] with_query [, ...] ]`. The bottom of the synopsis section notes: "where `from_item` can be one of:".

```
[ WITH [ RECURSIVE ] with_query [, ...] ]
SELECT [ ALL | DISTINCT [ ON ( expression [, ...] ) ] ]
      [ * | expression [ AS output_name ] [, ...] ]
      [ FROM from_item [, ...] ]
      [ WHERE condition ]
      [ GROUP BY [ ALL | DISTINCT ] grouping_element [, ...] ]
      [ HAVING condition ]
      [ WINDOW window_name AS ( window_definition ) [, ...] ]
      [ { UNION | INTERSECT | EXCEPT } [ ALL | DISTINCT ] select ]
      [ ORDER BY expression [ ASC | DESC | USING operator ] [ NULLS { FIRST | LAST } ] [, ...] ]
      [ LIMIT { count | ALL } ]
      [ OFFSET start [ ROW | ROWS ] ]
      [ FETCH { FIRST | NEXT } [ count ] { ROW | ROWS } { ONLY | WITH TIES } ]
      [ FOR { UPDATE | NO KEY UPDATE | SHARE | KEY SHARE } [ OF table_name [, ...] ] [ NOWAIT | SKIP LOCKED ] [ ... ] ]
```



# 9 Next Lecture

- Normalization
- Functional Dependencies
- Normal Forms
  - 1NF, 2NF, 3NF, BCNF, 4NF, 5NF, 6NF

